

# The Lagrangian Subspace Noise Problem: A New Framework for Post-Quantum Cryptography

Kwanghoo Choo\*

June 2026

## Abstract

We introduce the Lagrangian Subspace Noise (LSN) problem — a variant of the Learning Parity with Noise (LPN) problem where the secret is a Lagrangian subspace of a symplectic vector space rather than a linear function. The secret space, the Lagrangian Grassmannian  $\text{Lagr}(2n, \mathbb{F}_2)$ , has cardinality  $2^{n(n+1)/2+O(1)}$ , exponentially larger than LPN's  $2^n$  yet tightly constrained by the symplectic form. We prove an unconditional exponential Statistical Query (SQ) lower bound of  $\Omega(2^n)$  queries via an explicit symplectic spread, and a conditional  $2^{2n-O(1)}$  bound under a precisely-stated extremality conjecture, together with an exact pairwise-correlation formula and no asymptotic error term. We establish that linear and low-degree polynomial feature-map reductions are information-theoretically impossible, that adaptive linear SQ algorithms have only exponentially small advantage, and we map the exact standard-model gap that remains for general non-linear reductions. We construct two cryptographic primitives from LSN: a succinct non-interactive argument (LSN-SNARK) with  $O(n^2)$  R1CS constraints ( $\approx 4,225$  for  $n = 65$ , compared to millions for ZK-lattice proofs), and a key-encapsulation mechanism (LSN-KEM) with concatenated polar-code reconciliation achieving 80-bit security with public keys of 1.78 KB and ciphertexts of 288 B. All security claims are explicitly classified as theorem, evidence, or conjecture, and the remaining open problems are stated precisely.

**Keywords.** Post-quantum cryptography, Learning Parity with Noise, statistical query model, symplectic geometry, polar codes, zero-knowledge proofs.

## Acknowledgements

During the preparation of this work, the author used large language models (Claude, Kimi, and Codex) to assist with literature review, mathematical exploration, code prototyping, and manuscript drafting. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the accuracy, integrity, and final content of the work.

**Data and code availability.** All Python prototypes, test vectors, and the L<sup>A</sup>T<sub>E</sub>X source are available at <https://github.com/Gattochoo/lsn-pq>.

**Conflict of interest.** The author declares no competing interests.

---

\*ORCID: 0009-0005-5682-8098. Email: [gattochoo@gmail.com](mailto:gattochoo@gmail.com)

# 1 Introduction

## 1.1 Background

The Learning Parity with Noise (LPN) problem, introduced by Blum *et al.* [BFKL94], asks to recover a secret linear function  $s \in \mathbb{F}_2^k$  given noisy inner products  $(a_i, \langle a_i, s \rangle \oplus e_i)$ . Despite decades of cryptanalytic effort, no polynomial-time algorithm is known for LPN with constant noise rate. The best classical attacks run in sub-exponential time  $2^{O(n/\log n)}$  using the Blum–Kalai–Wasserman (BKW) algorithm [BKW03] and its modern refinements [LF06, Lyu05, EKM17, GJL14, ZJW16, WS21]. Grover’s algorithm reduces the naive key-search cost from  $2^k$  to  $2^{k/2}$  [Gro96], and quantum variants of BKW remain the dominant quantum threat; no published exponential quantum speed-up over the best classical methods is known [Kir11].

LPN has been the foundation for a wide range of cryptographic primitives, including OWFs, PKE, OT, and MPC [ABW10, Ale11, DGM14, BCGI18]. Its main theoretical limitation is the absence of a worst-case hardness reduction comparable to Regev’s LWE reduction from lattice problems [Reg05]. LPN hardness is therefore treated as a *primitive assumption*, accepted on the basis of sustained cryptanalytic resistance rather than reducibility.

Post-quantum cryptography currently rests on six classical hardness families: multivariate quadratic equations (MQ), lattices (LWE/SIS), error-correcting codes (LPN/syndrome decoding), hash functions, Ising models, and tensor problems. A seventh family — genuinely new, structurally distinct, and not reducible to the existing six — has been sought since the inception of the field. LSN is the strongest current candidate for such a family.

## 1.2 The LSN Problem

We propose *Lagrangian Subspace Noise* (LSN), also referred to as *Learning Stabilizers with Noise* or *symplectic LPN*. The secret is not a linear function but a Lagrangian subspace  $L \subset \mathbb{F}_2^{2n}$  of a symplectic vector space. The learner receives samples

$$(a_i, b_i) \quad \text{where} \quad b_i = \mathbf{1}_L(a_i) \oplus e_i,$$

with  $a_i \sim \mathbb{F}_2^{2n}$  uniform,  $e_i \sim \text{Bernoulli}(p)$ , and  $\mathbf{1}_L$  the indicator of  $L$ . The secret space is the Lagrangian Grassmannian:

$$|\text{Lagr}(2n, \mathbb{F}_2)| = \prod_{i=1}^n (2^i + 1) = 2^{n(n+1)/2 + O(1)}. \quad (1)$$

This is exponentially larger than LPN’s  $2^n$  secrets, yet it is not an unstructured space: every secret satisfies the strong isotropy constraint  $L = L^{\perp\Omega}$ . It is exactly this tension — exponential size versus rigid structure — that creates both the hardness and the technical difficulty of LSN.

Recent independent work by Khesin *et al.* [KLP+25], Poremba–Quek–Shor [PQS26], and Lu *et al.* [LPQR26] has shown that LSN is at least as hard as constant-rate LPN even for a single logical qubit, and has established strong barriers against reducing symplectic LPN to standard LPN (for linear reductions, proven in the constant-rate sample regime  $m = \Theta(n)$ ; see Section 9). These results position LSN as the strongest current candidate for a genuinely new post-quantum hardness family outside the six classical ones (multivariate-quadratic, lattice, code, hash, Ising, tensor).

## 1.3 Contributions

Our contributions are organized around proof, barriers, and constructions.

1. **Exact SQ lower bound (Section 6).** We prove an unconditional  $\Omega(2^n)$  SQ query lower bound via an explicit symplectic spread, and a conditional  $q \geq 2^{2n-O(1)}$  bound under a pencil-extremality conjecture. Both rest on an exact pairwise-correlation formula

$$\langle D_L, D_{L'} \rangle = \frac{(1-2p)^2}{p(1-p)} \cdot 2^{j-2n},$$

where  $j = \dim(L \cap L')$  and  $\langle D_L, D_{L'} \rangle$  denotes the pairwise correlation. For  $p = 1/4$  the coefficient is exactly  $4/3$ . A Markov concentration argument shows that a subset of size  $2^{2n}$  with average correlation  $O(2^{-2n})$  exists; the stronger worst-case bound remains open. The results hold for adaptive SQ algorithms.

2. **Reduction barriers (Section 9).** We prove that linear feature maps give zero advantage against LSN; polynomial feature maps of degree  $< n$  have error at least  $2^{-n}$ ; and adaptive linear SQ algorithms have advantage at most  $(1-2p)2^{-n}$  at any noise rate  $p$ . Together these barriers eliminate the natural reduction routes and localize the remaining open question to genuinely non-linear, adaptive reductions.

3. **Cryptographic primitives (Section 8).**

- **LSN-SNARK:** A preprocessing SNARK for LSN membership with  $O(n^2)$  R1CS constraints via the Siegel-chart (graph-Lagrangian) representation. Verified counts: 64 for  $n = 8$ , 1,764 for  $n = 42$ , 4,225 for  $n = 65$ .
- **LSN-KEM:** A CCA-secure KEM via concatenated polar codes. Direct polar coding at  $p = 1/4$  fails ( $Z_0 = 0.866$ , SC bound  $P_e > 1$ ), so we concatenate an inner repetition code with an outer polar code. For  $r = 7$  repetition we obtain effective noise  $p' = 0.0706$ , SC bound  $P_e < 2^{-81}$ , and parameters PK = 1.78 KB, CT = 288 B at 80-bit security. For  $r = 11$  we obtain  $P_e < 2^{-149}$  at 128-bit security.

4. **Rigorous caveats (Section 10).** Every claim is classified as theorem, evidence, or conjecture. We state explicitly the standard-model gap — the absence of a non-vacuous LPN(low-noise)  $\rightarrow$  LSN or LWE  $\rightarrow$  LSN reduction — that must close for full community acceptance.

## 2 Preliminaries

### 2.1 Symplectic Vector Spaces

Let  $V = \mathbb{F}_2^{2n}$  with standard symplectic form

$$\Omega(x, y) = \sum_{i=1}^n (x_{2i-1}y_{2i} + x_{2i}y_{2i-1}).$$

A subspace  $L \subset V$  is *isotropic* if  $\Omega|_{L \times L} = 0$ , and *Lagrangian* if it is isotropic of dimension  $n$ ; equivalently  $L = L^{\perp\Omega}$ . We write  $\text{Lagr}(2n, \mathbb{F}_2)$  for the Lagrangian Grassmannian.

**Proposition 2.1** (Lagrangian count [Mac71]).  $|\text{Lagr}(2n, \mathbb{F}_2)| = \prod_{i=1}^n (2^i + 1) = 2^{n(n+1)/2+O(1)}$ .

The intersection of two Lagrangians has even codimension. For  $L, L' \in \text{Lagr}(2n)$  let  $j = \dim(L \cap L')$ . The distribution of  $j$  is governed by the  $q$ -binomial coefficients at  $q = 2$ :

$$\Pr[j = k] = \frac{\begin{bmatrix} n \\ k \end{bmatrix}_2 \cdot 2^{(n-k)(n-k+1)/2}}{|\text{Lagr}(2n, \mathbb{F}_2)|}. \quad (2)$$

## 2.2 Fourier Transform on the Symplectic Space

The unnormalized  $\Omega$ -Fourier transform of  $f: \mathbb{F}_2^{2n} \rightarrow \mathbb{R}$  is

$$\mathcal{F}_\Omega[f](\xi) = \sum_{x \in \mathbb{F}_2^{2n}} f(x) (-1)^{\Omega(x, \xi)}.$$

**Lemma 2.2** (Self-duality). *For every  $L \in \text{Lagr}(2n, \mathbb{F}_2)$ ,*

$$\mathcal{F}_\Omega[\mathbf{1}_L] = 2^n \cdot \mathbf{1}_L.$$

*Proof.* If  $\xi \in L$ , then  $\Omega(x, \xi) = 0$  for all  $x \in L$ , so the sum is  $|L| = 2^n$ . If  $\xi \notin L$ , choose  $x_0 \in L$  with  $\Omega(x_0, \xi) = 1$ ; the involution  $x \mapsto x + x_0$  pairs terms of opposite sign, so the sum is 0.  $\square$

## 2.3 Statistical Query Model

Let  $\mathcal{D}$  be a class of distributions and  $D_0$  a reference distribution. The *statistical dimension with average correlation* [FGR+17] is

$$\text{SDA}(B(\mathcal{D}, D_0), \gamma) = \max \left\{ d : \forall S \subseteq \mathcal{D}, |S| \geq \frac{|\mathcal{D}|}{d} \Rightarrow \frac{1}{|S|^2} \sum_{D, D' \in S} |\langle D, D' \rangle| \leq \gamma \right\}.$$

**Theorem 2.3** (Feldman *et al.* [FGR+17]). *If  $\text{SDA}(B(\mathcal{D}, D_0), \gamma) = d$ , any SQ algorithm distinguishing  $\mathcal{D}$  from  $D_0$  with success probability  $\alpha > 1/2$  requires  $q \geq (2\alpha - 1)d$  queries to  $\text{VSTAT}(1/(3\gamma))$ .*

The theorem applies to randomized *adaptive* SQ algorithms, so our lower bound will inherit adaptivity automatically.

## 2.4 Notation

We write  $D_L$  for the LSN distribution with secret  $L$  and noise rate  $p$ , and  $D_0$  for the noise-only distribution in which  $b \sim \text{Bernoulli}(p)$  independently of  $a$ . All logarithms are base 2 unless stated otherwise.

## 3 Related Work

**LPN foundations and cryptanalysis.** LPN was introduced by Blum, Furst, Kearns, and Lipton [BFKL94] as a foundation for cryptographic primitives. The dominant classical attack is BKW [BKW03], later optimized with fast Walsh–Hadamard transforms [LF06], covering codes [GJL14, ZJW16], and hybrid information-set-decoding techniques [EKM17, WS21]. Lyubashevsky [Lyu05] gave a polynomial-sample variant, and a comprehensive unified attack framework was developed by Esser, Kübler, and May [EKM17]. Quantum speed-ups are limited to Grover [Gro96] and quantum BKW [Kir11]; no published exponential improvement is known.

**LPN variants.** Many structured variants have been proposed to improve efficiency: Ring-LPN [HKL+12], LPN with structured noise [DP12], and the Learning with Rounding (LWR) problem [BPR12]. All of these retain a secret space of size  $2^{O(n)}$ . LSN differs fundamentally in that its secret space has size  $2^{n^2/2+O(n)}$  while remaining polynomial-time sampleable via symplectic linear algebra.

**Symplectic and quantum coding.** Stabilizer codes, the quantum analogue of classical linear codes, are naturally described by Lagrangian subspaces of  $\mathbb{F}_2^{2n}$  [Got97, CRSS98, NC00]. Quantum state tomography of stabilizer states under depolarizing noise is a closely related learning problem [GL22]. The recent independent work of Khesin *et al.* [KLP+25] proved that constant-rate LPN reduces to LSN even for  $k = 1$ , while Lu *et al.* [LPQR26] proved that linear reductions cannot reduce sympLPN to LPN in the constant-rate sample regime  $m = \Theta(n)$ ; for  $m = \omega(n)$  their bound shows error weight  $1/2 - \delta$ , which they note is not sufficient to fully rule out decodability (though they expect the barrier to extend to all  $m = \text{poly}(n)$ ).

**Statistical query hardness.** The SQ model was introduced by Kearns [Kea98] and developed by Blum *et al.* [BFJ+94]. Feldman, Grigorescu, Reyzin, and Vempala [FGR+17] gave the general SDA framework we use, and Diakonikolas *et al.* [DKS17] applied it to high-dimensional robust estimation. Quantum algorithms that access only classical samples can achieve at most a quadratic Grover speed-up over classical algorithms. Since the SQ lower bound is exponential in  $n$ , such a speed-up does not affect the asymptotic security parameter, making SQ lower bounds relevant evidence for quantum hardness.

**Code-based cryptography and polar codes.** Code-based cryptography dates to McEliece [McE78] and Prange [Pra62]. Modern standards include Classic McEliece [ABD+21], HQC [AAC+22], and BIKE. Polar codes, introduced by Arikan [Ari09], are capacity-achieving with explicit construction; list decoding was introduced by Tal and Vardy [TV15], and finite-length analysis was developed by Mori–Tanaka [MT09], Trifonov [Tri12], and Hassani *et al.* [HVV14]. We use Bhattacharyya bounds for conservative parameter selection [Sei15].

**Post-quantum KEMs.** The NIST PQC standardization process selected Kyber (lattice) [SAB+22], Classic McEliece, and HQC [AAC+22] as KEM candidates. LSN-KEM offers competitive sizes with a fundamentally different assumption.

## 4 The LSN Problem

**Definition 4.1** (Membership- $\text{LSN}_{n,m,p}$ ). Let  $L \subset \mathbb{F}_2^{2n}$  be uniformly random Lagrangian. The LSN distribution  $D_L$  over  $\mathbb{F}_2^{2n} \times \mathbb{F}_2$  is generated by sampling  $a \sim \mathbb{F}_2^{2n}$ ,  $e \sim \text{Bernoulli}(p)$ , and outputting  $(a, b)$  where  $b = \mathbf{1}_L(a) \oplus e$ . Given  $m$  independent samples from  $D_L$ , recover  $L$ .

**Search vs. decision.** All hardness results in this paper are stated for the *decisional* problem: distinguishing  $D_L$  from the noise-only distribution  $D_0$  (where  $b \sim \text{Bernoulli}(p)$  independently of  $a$ ). The search-to-decision reduction follows the same template as for LPN [BFKL94]: given a decision oracle and a candidate Lagrangian  $L'$ , one can verify whether  $L' = L$  by testing whether the samples are consistent with  $D_{L'}$ ; a standard hybrid argument then yields search hardness from decision hardness. We therefore treat decisional LSN as the primitive assumption.

### 4.1 Two Formulations

Definition 4.1 is the *membership* (oracle) formulation: the query point  $a$  is freshly random for every sample, and the adversary sees it only together with its label. This is the natural form for statistical-query and information-theoretic analysis.

Cryptographic constructions and external hardness results use the *batch* formulation, where all query points are fixed in advance.

**Definition 4.2** (Batch- $\text{LSN}_{n,m,p}$ ). Let  $L \subset \mathbb{F}_2^{2n}$  be uniformly random Lagrangian, and let  $A = (a_1, \dots, a_m) \in (\mathbb{F}_2^{2n})^m$  be a fixed sequence of query points (public parameters). The batch

distribution  $D_L^{(A)}$  outputs labels  $b = (b_1, \dots, b_m)$  where  $b_j = \mathbf{1}_L(a_j) \oplus e_j$  with independent  $e_j \sim \text{Bernoulli}(p)$ . Given  $A$  and  $b$ , recover  $L$ .

The columns of  $A$  may be uniform random, pseudorandom (as in our KEM, Section 8.2), or structured.

**Definition 4.3** (sympLPN $_{n,p}$  [KLP+25, LPQR26]). Let  $A \in \mathbb{F}_2^{2n \times n}$  be a public matrix whose columns are *isotropic*:  $S_A := A^\top \Omega A = 0$  (every pair of columns is orthogonal under the symplectic form). Let  $x \in \mathbb{F}_2^n$  be a secret vector,  $e \sim \text{Bernoulli}(p)^{2n}$  a noise vector, and  $y = Ax + e \in \mathbb{F}_2^{2n}$ . Given  $(A, y)$ , recover  $x$ .

### Relationship.

1. **Batch-LSN with uniform  $A \equiv$  membership-LSN.** When the columns of  $A$  are uniform and independent, the adversary's view  $(A, b)$  in batch-LSN is identical in distribution to  $m$  independent samples from the membership oracle. Hence SQ lower bounds and information-theoretic floors proved for membership-LSN transfer *verbatim* to batch-LSN with uniform  $A$ .
2. **Pseudorandom  $A$  (KEM).** Our KEM (Section 8.2) uses a pseudorandom  $A$  generated from a short seed. Security reduces to batch-LSN with uniform  $A$  plus one PRG step.
3. **sympLPN is a different problem.** True sympLPN (Definition 4.3) has linear labels  $y = Ax + e$  and secret  $x$ ; the isotropic constraint  $S_A = 0$  is a public quadratic relation on the public parameters. It is not equivalent to batch-LSN (membership labels) without a pinned bridge theorem. LPQR26 study sympLPN; our transport theorems (Section 9) use only the  $S_A = 0$  structure on the public matrix and therefore apply to any formulation with isotropic columns, regardless of the label structure.

Throughout the paper we state which formulation each theorem addresses; unless noted otherwise, “LSN” refers to the membership form.

## 4.2 Information-Theoretic Floors

The per-sample information content of membership-LSN is exponentially small, which makes the membership form statistically trivial at polynomial sample complexity and *re-points* the hardness question to the batch formulation.

**Proposition 4.4** (Per-sample mutual information). *For membership-LSN with noise rate  $p$ ,*

$$I(L; (a, b)) = \Theta(2^{-n}).$$

*Proof.* Given  $L$ , the label is  $b = \mathbf{1}_L(a) \oplus e$  with  $e \sim \text{Bernoulli}(p)$ , so  $H(b | a, L) = H_2(p)$  exactly. For  $a \neq 0$ , a uniformly random  $a$  lies in  $L$  with probability  $(2^n - 1)/(2^{2n} - 1) = 2^{-n} + O(2^{-2n})$ , hence  $\Pr[b = 1 | a] = p + 2^{-n}(1 - 2p) + O(2^{-2n})$ ; for  $a = 0$  we have  $\Pr[b = 1 | a = 0] = 1 - p$ , and its entropy is  $H_2(1 - p) = H_2(p)$ , but this occurs with probability  $2^{-2n}$  and contributes negligibly. Thus  $H(b | a) = H_2(p + 2^{-n}(1 - 2p)) + O(2^{-2n})$  and the mutual information is  $\Theta(2^{-n})$  by the Taylor expansion  $H_2(p + \varepsilon) = H_2(p) + \varepsilon \log_2 \frac{1-p}{p} + O(\varepsilon^2)$ .  $\square$

**Proposition 4.5** ( $\chi^2$  divergence and sample complexity). *For any fixed Lagrangian  $L$ , the  $\chi^2$  divergence from  $D_0$  is*

$$\chi^2(D_L \| D_0) = \frac{(1 - 2p)^2}{p(1 - p)} \cdot 2^{-n}.$$

*Consequently any distinguisher (even computationally unbounded) needs  $m = \Omega(2^n)$  samples to tell  $D_L$  from  $D_0$  with constant advantage, and any search algorithm needs  $m = \Omega(n^2 \cdot 2^n)$  samples to recover  $L$  with constant probability.*

*Proof.* From the likelihood-ratio calculation in Appendix A, the squared likelihood ratio has expectation  $\mathbb{E}_{D_0}[\ell_L^2] = \frac{(1-2p)^2}{p(1-p)} \cdot 2^{-n}$ . Since  $\mathbb{E}_{D_0}[\ell_L] = 0$ , this is exactly  $\chi^2(D_L \| D_0)$ . The standard  $\chi^2$ -based sample-complexity bound (Le Cam’s method) gives  $m = \Omega(1/\chi^2) = \Omega(2^n)$  for decision. For search, Fano’s inequality with  $|\text{Lagr}| = 2^{n(n+1)/2+O(1)}$  and per-sample information  $\Theta(2^{-n})$  yields  $m = \Omega(n^2 \cdot 2^n)$ .  $\square$

**Consequence for reductions.** Propositions 4.4 and 4.5 imply that *membership-LSN with poly( $n$ ) samples is information-theoretically secure*: no reduction, linear or non-linear, adaptive or static, can recover  $L$  or distinguish  $D_L$  from  $D_0$  with only polynomially many samples. This makes “membership-LSN  $\nrightarrow$  LPN with poly( $n$ ) samples” trivially true and cryptographically empty. The hardness assumption that backs our KEM (Section 8.2) is *batch-LSN with pseudorandom  $A$*  (Definition 4.2). The barrier studied by LPQR26 is *sympLPN* (Definition 4.3), a decoding problem with linear labels and public isotropic matrix.

### 4.3 Distance Distribution and Expected Intersection

The correlation between two LSN distributions depends only on the dimension  $j = \dim(L \cap L')$  of the intersection of their secrets. The distribution of  $j$  is:

**Theorem 4.6** (Distance distribution). *For uniform  $L, L' \in \text{Lagr}(2n, \mathbb{F}_2)$ ,*

$$\Pr[j = k] = \frac{\binom{n}{k}_2 \cdot 2^{(n-k)(n-k+1)/2}}{|\text{Lagr}(2n, \mathbb{F}_2)|}.$$

*Numerical evaluation shows that  $\mathbb{E}[j] \rightarrow 0.76$  and  $\text{Var}(j) \rightarrow 0.60$  as  $n \rightarrow \infty$ .*

The rapid decay of Equation (2) means that two random Lagrangians typically intersect in dimension 0 or 1; intersections of dimension  $\geq 14$  occur with probability  $< 2^{-100}$  for  $n \geq 42$ .

### 4.4 The Diluted-LPN View

In the Siegel chart  $L = \{(u, Mu) : u \in \mathbb{F}_2^n\}$  with  $M \in \text{Sym}(n, \mathbb{F}_2)$ , a sample  $(a, b)$  with  $a = (u, v)$  and  $b = 1$  reveals  $n$  linear equations  $v = Mu$  in the  $n(n+1)/2$  unknown entries of  $M$ . The catch is that true positives are drowned in noise.

**Proposition 4.7** (Dilution of true positives). *Conditioned on the label  $b = 1$ , the posterior probability that the query point lies in  $L$  is*

$$\Pr[a \in L \mid b = 1] = \frac{(1-p) 2^{-n}}{p + (1-2p) 2^{-n}} = \Theta(2^{-n}).$$

*At  $p = 1/4$  the exact constant is 3 asymptotically ( $\approx 3 \cdot 2^{-n}$  for large  $n$ ); for  $n = 3$  it is exactly  $3/10$ .*

*Proof.* Bayes’ rule:  $\Pr[a \in L \mid b = 1] = \Pr[b = 1 \mid a \in L] \Pr[a \in L] / \Pr[b = 1] = (1-p) \cdot 2^{-n} / (p + (1-2p) 2^{-n})$ .  $\square$

Proposition 4.7 reframes batch-LSN search as an  $\varepsilon$ -diluted LPN problem over  $\text{Sym}(n, \mathbb{F}_2)$ : with probability  $\varepsilon \approx 3 \cdot 2^{-n}$  a sample yields  $n$  clean linear equations in  $M$ ; otherwise it yields noise. This is LPN with secret dimension  $n(n+1)/2$  at an extreme dilution rate. The computational question becomes:

**Enrichment question.** Can any efficient selection rule (using past labeled samples and the symplectic structure) produce query points with  $\Pr[a \in L] \geq (1 + \delta) 2^{-n}$  for non-negligible  $\delta$ ? In other words, can the  $2^{-n}$  dilution be beaten?

A negative answer would be the cleanest quantitative statement of what  $\text{LSN} \setminus \text{LPN}$  is; a positive answer would yield a concrete attack improvement (and a step toward a 6.5th-family demotion). The SQ lower bounds of Section 6 suggest that enrichment is impossible for SQ-implementable selection rules: any such rule is a linear combination of SQ queries, and the correlation bounds limit its advantage to  $2^{-n}$ . A formal reduction from enrichment to distinguishing would require a short lemma (omitted here).

## 4.5 Parameters

Table 1 gives the exact security parameters derived in Section 6.

Table 1: LSN security parameters for  $p = 1/4$ .

| Security | $n$ | $\log_2(q_{\min})^{(a)}$ | $ \text{Lagr} $ | PK (KEM) |
|----------|-----|--------------------------|-----------------|----------|
| 80-bit   | 41  | 80.6                     | $\sim 2^{861}$  | 1.78 KB  |
| 128-bit  | 65  | 128.6                    | $\sim 2^{2145}$ | 2.78 KB  |
| 192-bit  | 97  | 192.6                    | $\sim 2^{4753}$ | —        |
| 256-bit  | 129 | 256.6                    | $\sim 2^{8385}$ | —        |

<sup>(a)</sup> The  $q_{\min}$  figures use the conditional bound of

Theorem 6.6; the unconditional spread bound (Theorem 6.4) gives  $\Omega(2^n)$ .

The  $q_{\min}$  figures use the exact pairwise-correlation formula of Lemma 6.1, so the 80-bit line is met *exactly* at  $n = 41$  with no asymptotic slack. The constant factor in the exponent is  $\approx 0.67$  (below 1), so parameters with  $2n < \lambda$  are not advised without a rigorous upper bound. PK sizes are shown only for the two KEM parameter sets derived in Section 8.2; 192/256-bit KEM parameters depend on the inner repetition length  $r$  and are left as standard extrapolations.

## 5 Decoder Landscape

We review the natural decoder families and explain why each fails against LSN at constant noise.

**Linear decoders.** A real-linear query has the form  $q(a, b) = \langle w, a \rangle + c \cdot b$ . Its expectation under  $D_L$  is

$$\mathbb{E}_{D_L}[q] = c \cdot \Pr[b = 1] = c \cdot (p + 2^{-n}(1 - 2p)),$$

which is independent of  $L$ . Hence real-linear decoders have exactly zero  $L$ -identification power.  $F_2$ -linear queries can achieve advantage at most  $(1 - 2p)2^{-n}$  (Theorem 6.8).

**Polynomial decoders.** The indicator  $\mathbf{1}_L$  has an exact degree- $n$  polynomial representation

$$\mathbf{1}_L(x) = \prod_{j=1}^n (1 + \langle w_j, x \rangle)$$

where  $\{w_j\}$  is a basis for the annihilator of  $L$ . Any polynomial of degree  $< n$  misses at least one factor and therefore has error  $\geq 2^{-n}$ .

**List decoding.** The list size needed to enumerate all candidate Lagrangians is  $|\text{Lagr}| = 2^{n^2/2 + O(n)}$ . Even with pruning, list decoding is infeasible for  $n \geq 41$ .



**BKW and ISD.** BKW would require treating the secret as an  $n^2$ -bit object, giving complexity  $2^{\Omega(n^2/\log n)}$ . Information-set decoding does not apply directly because the code ensemble (Lagrangian indicators) is highly structured rather than random linear.

**Structural attacks.** For very low noise ( $p = o(1/n)$ ), the span of the positive samples reveals  $L$ . At constant noise  $p = 1/4$ , [LPQR26] and our experiments show that structural attacks are ineffective: the span of positives has full dimension after  $O(n)$  samples.

**Empirical sample-complexity lower bound.** We ran brute-force experiments for  $n = 3, 4, 5$  using transvection-orbit enumeration of the full Lagrangian family. In each trial we sampled  $L \leftarrow \text{Lagr}$  uniformly, drew  $m$  noisy samples from  $\text{LSN}_{L,1/4}$ , and let a brute-force decoder enumerate all Lagrangians, picking the one with the highest agreement. The recovery success rates (200 trials for  $n = 3$ , 50 for  $n = 4$ , 20 for  $n = 5$ ) were:

| $n$ | $ \text{Lagr} $ | $m = 2^{2n-2}$ | $m = 2^{2n-1}$ | $m = 2^{2n}$ | $m = 2 \cdot 2^{2n}$ |
|-----|-----------------|----------------|----------------|--------------|----------------------|
| 3   | 135             | 11.5%          | 18.5%          | 46.5%        | 80.0%                |
| 4   | 2,295           | 6.0%           | 24.0%          | 62.0%        | 98.0%                |
| 5   | 75,735          | 5.0%           | 20.0%          | 75.0%        | 100% (20/20)         |

Two trends are visible. First, the success rate at the floor  $m = 2^{2n}$  *increases* with  $n$  (46.5%, 62%, 75%), consistent with  $2^{2n}$  being the asymptotic threshold (the constant factor in front appears to be below 1 and to matter less as  $n$  grows). Second, halving the budget to  $m = 2^{2n-1}$  already drops recovery below 25% even at  $n = 5$ , so the threshold sits at the  $2^{2n}$  scale rather than a constant factor below it. These experiments *empirically support* a  $2^{2n}$  sample-complexity scale for the brute-force ML decoder, matching the pairwise floor of Lemma 6.1; they do not by themselves rule out a more sample-efficient decoder—the rigorous lower bound is the SQ query bound of Section 6.

## 6 Statistical Query Lower Bound

This section contains the main technical result: an exponential SQ lower bound with an exact correlation formula.

### 6.1 Exact Pairwise Correlation

Let  $D_0$  be the noise-only distribution. The likelihood ratio of  $D_L$  with respect to  $D_0$  is

$$\ell_L(a, b) = \frac{dD_L}{dD_0}(a, b) - 1 = \mathbf{1}_L(a) \cdot \beta \cdot \left( \mathbf{1}_{b=1} - \mathbf{1}_{b=0} \frac{p}{1-p} \right),$$

where  $\beta = (1 - 2p)/p$ .

**Lemma 6.1** (Exact pairwise correlation). *For  $L, L' \in \text{Lagr}(2n, \mathbb{F}_2)$  with  $j = \dim(L \cap L')$ ,*

$$\langle D_L, D_{L'} \rangle = \frac{(1 - 2p)^2}{p(1 - p)} \cdot 2^{j-2n}. \quad (3)$$

*Proof.* Because  $a$  and  $b$  are independent under  $D_0$ , the inner product factors:

$$\langle D_L, D_{L'} \rangle = \mathbb{E}_a[\mathbf{1}_L(a)\mathbf{1}_{L'}(a)] \cdot \mathbb{E}_b \left[ \beta^2 \left( \mathbf{1}_{b=1} - \mathbf{1}_{b=0} \frac{p}{1-p} \right)^2 \right].$$

The first factor is  $|L \cap L'|/2^{2n} = 2^{j-2n}$ . The second factor is

$$\beta^2 \left( p + (1-p) \frac{p^2}{(1-p)^2} \right) = \frac{(1-2p)^2}{p^2} \cdot \frac{p}{1-p} = \frac{(1-2p)^2}{p(1-p)}.$$

Multiplying yields Equation (3). For  $p = 1/4$  the coefficient is exactly 4/3.  $\square$

## 6.2 Average Correlation

Taking expectation over  $j$  gives the average pairwise correlation.

**Lemma 6.2** (Average correlation). *Let  $C_n = \mathbb{E}[2^j]$  computed from Theorem 4.6. Then  $C_n \rightarrow 2$  and*

$$\rho_{\text{avg}} = \frac{(1-2p)^2}{p(1-p)} \cdot C_n \cdot 2^{-2n}.$$

## 6.3 SDA Concentration

**Lemma 6.3** (Existence of low-correlation subset). *Let  $\gamma = 2\rho_{\text{avg}}$ . There exists a subset  $\mathcal{D}' \subset \{D_L\}$  of size  $|\mathcal{D}'| = 2^{2n}$  with average correlation at most  $\gamma$ .*

*Proof.* Consider a uniformly random subset  $S \subset \text{Lagr}(2n)$  of size  $M = 2^{2n}$ . By symmetry of the Lagrangian Grassmannian under  $\text{Sp}(2n)$ ,  $\mathbb{E}[\rho(S)] = \rho_{\text{avg}}$ . Markov's inequality gives

$$\Pr[\rho(S) > 2\rho_{\text{avg}}] < \frac{\mathbb{E}[\rho(S)]}{2\rho_{\text{avg}}} = \frac{1}{2}.$$

Hence with probability at least  $1/2$ , a random subset of size  $2^{2n}$  has average correlation at most  $\gamma = 2\rho_{\text{avg}}$ . This proves existence of such a subset. A worst-case bound over *all* subsets of size  $2^{2n}$  would require a structural argument; we discuss this in Section 10.  $\square$

## 6.4 Unconditional SQ Bound: The Symplectic Spread

We first give an unconditional exponential lower bound on a concrete worst-case promise problem. The bound is weaker in both query count and VSTAT strength than the conditional average-case bound of Theorem 6.6, but it requires no unproven structural hypothesis.

Let  $V = \mathbb{F}_{2^n} \times \mathbb{F}_{2^n}$  with symplectic form  $\omega((a, b), (c, d)) = \text{Tr}(ad) + \text{Tr}(bc)$ , where  $\text{Tr}$  is the trace from  $\mathbb{F}_{2^n}$  to  $\mathbb{F}_2$ . The *Desarguesian symplectic spread*  $\mathcal{S}^* = \{L_\lambda\}_{\lambda \in \mathbb{F}_{2^n}} \cup \{L_\infty\}$  consists of  $d_0 = 2^n + 1$  Lagrangians

$$L_\lambda = \{(x, \lambda x) : x \in \mathbb{F}_{2^n}\}, \quad L_\infty = \{0\} \times \mathbb{F}_{2^n},$$

which are pairwise transversal:  $\dim(L \cap L') = 0$  for all distinct  $L, L' \in \mathcal{S}^*$ .

**Theorem 6.4** (Unconditional spread bound). *Let  $\kappa = (1-2p)^2/(p(1-p))$  and fix  $1 \leq t \leq n-1$ . Set  $\bar{\gamma}_t = \kappa \cdot 2^{-(2n-t)}$ . Every subset  $T \subseteq \mathcal{S}^*$  with  $|T| \geq 2^{n-t}$  has diagonal-inclusive average correlation at most  $2\bar{\gamma}_t$ . Consequently  $\text{SDA}(\mathcal{S}^*, 2\bar{\gamma}_t) \geq 2^t$ , and any SQ algorithm that is correct on every  $L \in \mathcal{S}^*$  requires  $\Omega(2^t)$  queries to  $\text{VSTAT}(1/(6\bar{\gamma}_t))$ .*

At  $t = n-1$  this yields an unconditional  $\Omega(2^n)$  query lower bound at VSTAT strength  $O(2^n)$ .

*Proof.* For distinct  $L, L' \in \mathcal{S}^*$  we have  $j = \dim(L \cap L') = 0$ , so by Lemma 6.1 the pairwise correlation is exactly  $\langle D_L, D_{L'} \rangle = \kappa \cdot 2^{-2n} = \bar{\gamma}_t \cdot 2^{-t} \leq \bar{\gamma}_t$ . The diagonal terms contribute  $\beta = \kappa \cdot 2^{-n}$  each, so for  $|T| \geq 2^{n-t}$

$$\frac{1}{|T|^2} \sum_{D, D' \in T} |\langle D, D' \rangle| \leq \frac{|T|^2 \cdot \bar{\gamma}_t + |T| \cdot \beta}{|T|^2} \leq \bar{\gamma}_t + \frac{\beta}{2^{n-t}} \leq 2\bar{\gamma}_t.$$

Hence  $\text{SDA}(\mathcal{S}^*, 2\bar{\gamma}_t) \geq |\mathcal{S}^*|/2^{n-t} \geq 2^t$ . Applying Theorem 2.3 with  $\alpha = 2/3$  gives the query bound.  $\square$

**Honesty notes.** (i) **Worst-case promise only.** The spread has measure  $|\mathcal{S}^*|/|\text{Lagr}(2n, \mathbb{F}_2)| \sim 2^n/2^{n(n+1)/2} \rightarrow 0$ , so Theorem 6.4 does *not* cover the average-case (uniform  $L$ ) problem; it protects only the promise problem “ $L \in \mathcal{S}^*$ ?”. (ii) **VSTAT trade-off.** The unconditional bound is  $\Omega(2^n)$  queries at  $\text{VSTAT}(O(2^n))$ ; the conditional bound below is  $2^{2n-O(1)}$  queries at  $\text{VSTAT}(\Theta(2^{2n}))$ . The two parameters are not interchangeable. (iii) The subfamily restriction is valid for the promise problem: any algorithm solving LSN over the full family is correct on  $\mathcal{S}^*$  a fortiori.

## 6.5 Conditional SQ Bound

The stronger  $2^{2n}$ -scale bound rests on a structural conjecture about extremal subsets.

**Conjecture 6.5** (Pencil extremality). *There is an absolute constant  $c$  such that every subset  $\mathcal{D}' \subseteq \{D_L\}$  of size  $|\mathcal{D}'| \geq |\text{Lagr}(2n, \mathbb{F}_2)|/2^{2n-c}$  has average correlation at most  $5\rho_{\text{avg}}$ .*

Motivation:  $k = 2$  isotropic pencils force any threshold above  $4\rho_{\text{avg}}$  (their expected  $2^{\dim(L \cap L')}$  ratio tends to 4 from below), while  $k \geq 3$  pencils fall below the size threshold  $|\text{Lagr}|/2^{2n}$  for  $n \geq 4$ , and mixtures of pencils with random fill dilute quadratically. A proof would require classifying near-extremal subsets.

**Theorem 6.6** (Conditional main SQ lower bound). *Assume Conjecture 6.5. Any SQ algorithm distinguishing LSN (uniform  $L$ ) from  $D_0$  with probability  $> 2/3$  requires  $q \geq 2^{2n-O(1)}$  queries to  $\text{VSTAT}(1/(15\rho_{\text{avg}}))$ .*

*Proof.* Under Conjecture 6.5 with  $c = 0$ , every subset of size at least  $|\text{Lagr}(2n, \mathbb{F}_2)|/2^{2n}$  has average correlation at most  $5\rho_{\text{avg}}$ . Hence  $\text{SDA}(B(\{D_L\}, D_0), 5\rho_{\text{avg}}) \geq 2^{2n}$ . Applying Theorem 2.3 with  $\alpha = 2/3$  and  $\gamma = 5\rho_{\text{avg}}$  yields  $q \geq (4/3 - 1) \cdot 2^{2n} = 2^{2n-O(1)}$  queries to  $\text{VSTAT}(1/(15\rho_{\text{avg}}))$ .  $\square$

**Theorem 6.7** (Adaptive SQ). *Theorem 6.6 holds for randomized adaptive SQ algorithms.*

*Proof.* Theorem 2.3 applies to randomized adaptive SQ algorithms [FGR+17, Theorem 3.7], so the bound is inherited directly.  $\square$

## 6.6 Adaptive Linear SQ is Blocked

**Theorem 6.8** (Adaptive linear SQ barrier). *For any real-linear query  $q(a, b) = \langle w, a \rangle + c \cdot b$ , the expectation  $\mathbb{E}_{D_L}[q]$  is independent of  $L$ . For any  $F_2$ -linear query  $q(a, b) = \langle w, a \rangle \oplus c \cdot b$ , the maximum distinguishing advantage against  $D_0$  is  $(1 - 2p)2^{-n}$ ; at  $p = 1/4$  this is  $2^{-n-1}$ .*

*Proof.* For real-linear queries, direct computation gives  $\mathbb{E}_a[\langle w, a \rangle] = 0$  and  $\mathbb{E}_{D_L}[b] = p + 2^{-n}(1 - 2p)$ , which is independent of  $L$ . For an  $F_2$ -linear query  $q(a, b) = \langle w, a \rangle \oplus c \cdot b$ : if  $c = 0$  and  $w \neq 0$  then  $\mathbb{E}_{D_L}[q] = 1/2$  for every  $L$  (the constant query  $w = 0, c = 0$  is trivially  $L$ -independent); if  $c = 1$  and  $w = 0$ , then  $q = b$  and the advantage over  $D_0$  is exactly  $(1 - 2p)2^{-n}$ . If  $c = 1$  and  $w \neq 0$ , write  $q = b \oplus \langle w, a \rangle$ . Under  $D_L$  we have

$$\mathbb{E}_{D_L}[q] = \frac{1}{2} + (1 - 2p)(2^{-n} - 2\mathbb{E}_a[\langle w, a \rangle \cdot \mathbf{1}_L(a)]).$$

If  $w \in L^\perp$  then  $\mathbb{E}_a[\langle w, a \rangle \cdot \mathbf{1}_L(a)] = 0$  and the advantage is  $(1 - 2p)2^{-n}$ . If  $w \notin L^\perp$  then  $\mathbb{E}_a[\langle w, a \rangle \cdot \mathbf{1}_L(a)] = 2^{-n-1}$  and the expectation collapses to  $1/2$ , giving advantage 0. Thus no  $F_2$ -linear query has advantage larger than  $(1 - 2p)2^{-n}$ .  $\square$

## 7 Quantum Analysis

### 7.1 Grover Search

The secret space has size  $|\text{Lagr}(2n, \mathbb{F}_2)| = 2^{n^2/2+O(n)}$ . A Grover search over this space requires  $2^{n^2/4+O(n)}$  iterations, each costing at least  $\Omega(mn)$  to evaluate an LSN sample set. The total cost  $2^{n^2/4+O(n)}$  dominates all classical bounds for  $n \geq 4$ , so Grover is not a practical threat.

### 7.2 Quantum BKW

A quantum version of BKW can speed up the inner Walsh–Hadamard transform and the search for low-weight parity checks, yielding complexity  $2^{O(n^2/\log n)}$  [Kir11]. Since  $n^2/\log n \gg 2n$  for all  $n \geq 4$ , this exceeds the SQ lower bound and is irrelevant for security parameter selection.

### 7.3 Quantum SQ

Because LSN samples are classical bits, any quantum algorithm must either access them through classical queries or perform a quantum walk over the sample space. A Grover search over the  $2^{n^2/2}$ -sized secret space still requires  $2^{n^2/4}$  iterations (Section 7.1), far above the security parameter. Thus the exponential SQ lower bound provides evidence (not proof) that no efficient quantum algorithm solves LSN.

**Conjecture 7.1** (Quantum LSN). *Any quantum algorithm solving  $LSN_{n,m,p}$  with constant  $p$  requires  $\Omega(2^n)$  queries.*

### 7.4 Dihedral Hidden Subgroup

The Lagrangian Grassmannian is not a group; it lacks a transitive abelian structure. Consequently, no reduction to the dihedral hidden-subgroup problem or its generalizations is known.

### 7.5 What Is and Is Not Covered

We summarize the quantum attack landscape honestly.

**Blocked vectors.** The following quantum approaches have been analyzed and ruled out for LSN at constant noise:

- **Hidden Subgroup Problem.** The Lagrangian Grassmannian lacks a group structure with coset decomposition; no reduction to DHSP or abelian HSP is known.
- **Grover search.** Complexity  $2^{n^2/2+O(n)}$  exceeds all classical bounds.
- **Quantum BKW.** Complexity  $2^{O(n^2/\log n)}$  is super-exponential in the security parameter.

**Not covered.** We do *not* claim a quantum lower bound for LSN. Quantum hardness is treated as a **conjecture** (Conjecture 7.1), analogous to the status of LWE and LPN. A rigorous quantum query lower bound would require new techniques beyond the SQ model and is left as future work.

## 8 Cryptographic Primitives

### 8.1 LSN-SNARK

We construct a preprocessing SNARK for the relation “ $x \in L$ ” given a committed Lagrangian  $L$ . The construction exploits the *Siegel-chart* (graph-Lagrangian) representation to reduce the R1CS circuit from  $O(n^3)$  to  $O(n^2)$  constraints.

**Siegel-chart representation.** Fix a public symplectic matrix  $S \in \text{Sp}(2n, \mathbb{F}_2)$  (e.g., generated from a standard seed). Every “generic” Lagrangian can be written as

$$L_M = S \cdot \{(x, Mx) : x \in \mathbb{F}_2^n\},$$

where  $M \in \text{Sym}(n, \mathbb{F}_2)$  is a symmetric matrix. The set of graph Lagrangians has cardinality  $|\text{Sym}(n)| = 2^{n(n+1)/2}$ , which differs from the full Lagrangian family  $|\text{Lagr}| = \prod_{i=1}^n (2^i + 1)$  by only a constant factor ( $\approx 2.38$ ). Consequently, restricting to graph Lagrangians does not affect asymptotic security; the brute-force complexity remains  $2^{n^2/2+O(n)}$  and the SQ lower bounds (Theorems 6.4 and 6.6) apply asymptotically unchanged because correlation structure is preserved under symplectic change of coordinates.

**Binding  $M$  to the public key.** The LSN-KEM public key (section 8.2) contains  $m$  noisy samples  $(x_i, y_i)$  generated from the secret Lagrangian  $L_M$ . To prove knowledge of  $M$  in a SNARK, the prover must bind their witness to the public key. The public key is therefore augmented with a commitment  $c = \text{Hash}(M)$ , where **Hash** is a SNARK-friendly hash function (e.g., Poseidon). This introduces preimage-resistance of **Hash** as a second assumption; we size **Hash** ( $\geq 256$ -bit output, Poseidon) so that quantum preimage search ( $2^{128}$ ) is at least the LSN level, making LSN the bottleneck. The SNARK circuit then verifies two things: (1) commitment opening ( $c = \text{Hash}(M)$ ), and (2) membership ( $x \in L_M$ ). Hashing the  $n(n+1)/2$  bits of  $M$  with Poseidon adds only  $O(n)$  constraints—packing the bits into  $\approx 9$  field elements for  $n = 65$  costs roughly 300–500 R1CS constraints. The full circuit size is therefore  $n^2 + O(n)$ , still  $O(n^2)$  and still orders of magnitude below ZK-lattice circuits.

**Membership circuit.** The core membership sub-circuit verifies  $x \in L_M$  in two steps:

1. Apply the public inverse  $S^{-1}$  to  $x$  to obtain  $(a, b) = S^{-1}x$ . This is a linear operation over  $\mathbb{F}_2$ , hence costs **zero** R1CS constraints.
2. Check  $b = Ma$ . Each entry  $b_i = \sum_{j=1}^n M_{ij}a_j$  is a dot product of length  $n$  over  $\mathbb{F}_2$ . Treating  $a$  as a witness vector yields  $n$  AND gates per row, for  $n^2$  constraints total; the final sums are linear and absorbed into the R1CS output side. Symmetry of  $M$  is free: we encode only the upper-triangular  $n(n+1)/2$  entries.

**Further optimization.** If  $a$  is provided directly as a public input (derived from  $x$  by the verifier), each product  $M_{ij}a_j$  becomes a selector—free when  $a_j = 0$  and a copy when  $a_j = 1$ . The check  $b_i = \sum_j M_{ij}a_j$  then collapses to a single linear constraint per row, giving  $n$  **constraints** over  $\mathbb{F}_2$  R1CS, or  $n(n+3)/2$  constraints over a prime field with boolean witnesses for  $M$ .

Table 2 compares LSN-SNARK with other ZK-PQC approaches.

Table 2: R1CS constraints for LSN membership vs. ZK-PQC full primitives.

| Scheme                         | Constraints ( $n = 65$ )                        | Asymptotic                 |
|--------------------------------|-------------------------------------------------|----------------------------|
| <b>LSN-SNARK membership</b>    | <b><math>\approx 4,225</math></b>               | <b><math>O(n^2)</math></b> |
| <b>LSN-SNARK full (+ hash)</b> | <b><math>\approx 4,500\text{--}4,700</math></b> | <b><math>O(n^2)</math></b> |
| SHA-256 (per block)            | $\approx 25,000$                                | $O(1)$                     |
| AES-128                        | $\approx 16,000$                                | $O(1)$                     |
| ZK-Kyber (full KEM verif.)     | millions                                        | $O(n \log n)$              |
| ZK-SPHINCS+ (full sig. verif.) | millions                                        | $O(\lambda \log \lambda)$  |

**Witness structure.** The witness  $\mathbf{w}$  consists of:

- $n(n+1)/2$  bits for the upper-triangular entries of  $M \in \text{Sym}(n, \mathbb{F}_2)$ ;
- $n$  bits for the vector  $a$  (the first half of  $S^{-1}x$ , computed as a linear function of  $x$  but included in the witness for the dot-product checks);
- auxiliary opening data for the hash commitment (omitted for deterministic hashes, negligible otherwise).

Total witness size:  $n(n+1)/2 + O(n) = O(n^2)$  bits.

**Verified counts.** The  $n^2$  prediction is exact for the membership sub-circuit when  $a$  is a witness variable:  $n = 8$  yields 64 constraints,  $n = 41$  (80-bit) yields 1,681, and  $n = 65$  (128-bit) yields 4,225. With public  $a$ , the optimized membership counts are 8, 41, and 65 respectively over  $\mathbb{F}_2$ . The *full* circuit (membership + Poseidon commitment opening) at  $n = 65$  is  $\approx 4,500$ – $4,700$  constraints depending on hash parameterization.

**Why this matters.** The membership relation  $x \in L_M$  thus has an unusually compact ZK circuit— $\approx 4,225$  R1CS constraints at  $n = 65$  (or 65 with public  $a$ ), small relative to hashing one SHA-256 block ( $\approx 25,000$ ). The full primitive (membership + commitment opening) stays below  $\approx 5,000$  constraints—about  $5\times$  smaller than a single SHA-256 block ( $\approx 25,000$ ) and roughly  $1,000\times$  smaller than published ZK-Kyber circuits (millions). **Scope.** The table compares *core-relation* circuits; full signature/KEM verification includes additional binding overhead that is comparable across schemes. Our claim is not a record for the full signature, but that the *membership* sub-circuit—the inner relation being proved—is exceptionally compact. That compactness helps building blocks such as *verifiable PQ encryption* and *recursive SNARK aggregation*, where many membership checks dominate the cost.

**Batch verification.** The compactness enables dramatic amortization when proving membership for  $k$  distinct elements  $x^{(1)}, \dots, x^{(k)}$  under the same secret  $M$ . The commitment  $c = \text{Hash}(M)$  is verified **once**; each additional membership check adds only  $n^2$  constraints (or  $n$  with public  $a$ ). For batch size  $k$ :

$$\text{Total constraints} = \underbrace{O(n)}_{\text{hash opening}} + k \cdot \underbrace{O(n^2)}_{\text{membership}}.$$

In contrast, ZK-Kyber verification requires independent NTT evaluations per proof, offering no comparable amortization. Concrete example ( $n = 65, k = 1000$ ): LSN uses  $\approx 500 + 1000 \cdot 4,225 \approx 4.2 \times 10^6$  constraints, whereas ZK-Kyber at  $\approx 5 \times 10^6$  constraints per proof would need  $\approx 5 \times 10^9$ —a  $1,000\times$  gap.

## 8.2 LSN-KEM

### 8.2.1 Why concatenated codes

Direct polar coding for LSN at  $p = 1/4$  fails. The Bhattacharyya parameter of  $\text{BSC}(p)$  is  $Z_0 = 2\sqrt{p(1-p)} = 0.866$ , too close to 1 for the finite-length polar code  $N = 2048$  to achieve any useful reliability. We therefore concatenate:

- an inner repetition code of length  $r$ , reducing the effective channel to  $\text{BSC}(p')$  where

$$p' = \sum_{k=\lceil r/2 \rceil + 1}^r \binom{r}{k} p^k (1-p)^{r-k};$$

- an outer polar code of length  $N = 2048$  and dimension  $K = 256$  over  $\text{BSC}(p')$ , decoded with SCL of list size  $L = 8$ .

Table 3 gives the parameters.

Table 3: LSN-KEM parameters for concatenated polar reconciliation.

| Security | $n$ | $r$ | $p'$   | SC bound   | SCL estimate |
|----------|-----|-----|--------|------------|--------------|
| 80-bit   | 41  | 7   | 0.0706 | $2^{-81}$  | $2^{-89}$    |
| 128-bit  | 65  | 11  | 0.0343 | $2^{-149}$ | $2^{-157}$   |

### 8.2.2 Construction

We use the random-oracle model (ROM) for the security proof.

**KeyGen( $1^\lambda$ ):** Let  $S \in \text{Sp}(2n, \mathbb{F}_2)$  be a public symplectic matrix. Choose a symmetric matrix  $M \leftarrow \text{Sym}(n, \mathbb{F}_2)$  uniformly and define  $L = S \cdot \{(x, Mx) : x \in \mathbb{F}_2^n\}$ . Choose  $\text{seed} \leftarrow \{0, 1\}^\lambda$ . For  $i = 1, \dots, m$  (where  $m = N \cdot r$ ), compute  $x_i = \text{PRG}(\text{seed}, i)$  and  $y_i = \mathbf{1}_L(x_i) \oplus e_i$ . Output  $\text{pk} = (\text{seed}, y_1, \dots, y_m)$  and  $\text{sk} = M$ .

**Security remark.** Restricting to graph Lagrangians reduces the key-space size by a constant factor  $|\text{Sym}(n)|/|\text{Lagr}| \rightarrow \prod_{i \geq 1} (1 + 2^{-i})^{-1} \approx 1/2.38$ , which does not affect the asymptotic security parameter. The SQ lower bounds (Theorems 6.4 and 6.6) and the empirical brute-force barriers (Section 5) are proved for membership-LSN; batch-LSN with uniform  $A$  is identical in distribution to membership-LSN (Section 4.1), so these barriers transfer verbatim to the batch formulation used here (pseudorandom  $A$  adds one PRG-indistinguishability step). Correlation structure is preserved under symplectic change of coordinates, so the bounds apply asymptotically unchanged. As a bonus, the secret key is more compact:  $M$  requires only  $n(n+1)/2$  bits versus  $2n^2$  bits for a general basis representation.

**Encaps:** Choose randomness  $\rho = (s, u)$  with  $s \leftarrow \{0, 1\}^\lambda$  and  $u \leftarrow \{0, 1\}^K$ . Compute a Fisher–Yates permutation  $\pi = \text{Permutation}(\text{PRG}(s))$  over  $[m]$ , and let  $\text{indices} = \pi[:Nr]$ . Group the selected indices into  $N$  blocks of length  $r$  and compute  $v_j = \text{majority}(y \text{ in block } j)$  for each block. Compute  $c = \text{PolarEncode}(u)$ , set  $\text{ct} = (s, v \oplus c)$ , and  $K = \text{Hash}(s, u)$ .

**Decaps:** Recompute  $\pi$  and the blocks. For each block compute  $v'_j = \text{majority}(\mathbf{1}_L(x))$  using the secret key. Because the  $x$ 's are derived from the same public seed, the  $v'_j$  are deterministic; they need not equal the noisy  $v_j$  sent by the encapsulator, but the difference  $v_j \oplus v'_j$  is exactly the effective noise seen by the polar decoder. Compute  $c' = \text{syn} \oplus v'$ , decode  $u' = \text{PolarSCLDecode}(c')$ , and output  $K = \text{Hash}(s, u')$  (or  $\perp$  if decoding fails).

The permutation-based index sampling guarantees distinct indices, eliminating birthday collisions. Public-key size is  $\lambda + Nr$  bits; ciphertext size is  $\lambda + N$  bits.

**Implementation note.** A small-scale Python prototype of the SCL decoder exhibited a bug (BLER = 1.0 for  $N = 128, 256$ ), contradicting the conservative Bhattacharyya bound and indicating an implementation issue rather than a parameter problem. Subsequent validation with the verified `komm` SC decoder [komm] revealed a second bug in the Bhattacharyya parameter indexing: the original code concatenated all bad channels before all good channels, whereas `komm`'s natural-order polar code requires interleaving ( $z_{2i} = 2z_i - z_i^2$ ,  $z_{2i+1} = z_i^2$ ). After correcting the frozen-set construction, all tested configurations yielded BLER = 0: at  $p' = 0.0706$

for  $N = 128, 256, 512$  with  $K = N/8$ , and at  $p' = 0.0343$  for all tested lengths. The previous  $N = 256$  anomaly (BLER = 0.115) was entirely due to the indexing mismatch, not incomplete polarization. These short-length, 200-trial simulations are *preliminary*: they certify only  $\text{BLER} \lesssim 1/200$ , not the design point  $2^{-80}$ , and they do not test the design length  $N = 2048$ . The conservative Arikan–Bhattacharyya bound ( $2^{-80}$  at  $r = 7$ ,  $2^{-148}$  at  $r = 11$ ) remains the design basis; full  $N = 2048$  validation is deferred to the production Rust implementation.

### 8.2.3 Correctness

**Theorem 8.1** (Correctness). *For  $r = 7$ ,  $\Pr[K \neq K'] < 2^{-80}$ ; for  $r = 11$ ,  $\Pr[K \neq K'] < 2^{-128}$ .*

*Proof.* Inner majority vote converts BSC(1/4) into BSC( $p'$ ). The outer polar code Bhattacharyya bound gives  $P_e \leq \frac{1}{2} \sum_{i \in \mathcal{I}} Z(W_i)$ , where  $\mathcal{I}$  is the information set. Numerical evaluation of the bound yields the claimed values; SCL decoding with  $L = 8$  improves them by a small constant factor.  $\square$

### 8.2.4 IND-CPA Security

**Theorem 8.2** (IND-CPA). *LSN-KEM is IND-CPA secure in the random-oracle model under the LSN hardness assumption.*

*Proof sketch.* The proof uses three game hops.

1. **Game 0**  $\rightarrow$  **Game 1**: Replace PRG with a truly random function. Indistinguishability follows from PRG security.
2. **Game 1**  $\rightarrow$  **Game 2**: Replace the public-key labels  $(y_1, \dots, y_m)$  with uniform random bits. We construct a direct reduction  $\mathcal{B}$  from KEM IND-CPA to LSN:  $\mathcal{B}$  receives an LSN challenge  $(x_i, y_i^*)$ , programs  $\text{seed}$  so that  $\text{PRG}(\text{seed}, i) = x_i$ , and sets  $\text{pk} = (\text{seed}, y_1^*, \dots, y_m^*)$ . The KEM challenger computes  $\text{ct}^*$  using all  $m = Nr$  samples. If the LSN challenge is real, the adversary sees Game 1; if random, Game 2. Any advantage  $\varepsilon$  in distinguishing the games translates to advantage  $\varepsilon$  against LSN. There is no  $1/r$  security loss because the reduction uses the LSN challenge directly.
3. **Game 2**  $\rightarrow$  **Game 3**: In Game 2 the vector  $v$  is uniform and independent of the polar codeword  $c = \text{PolarEncode}(u)$ . Hence  $\text{syn} = v \oplus c$  is uniform (one-time pad), and  $K = \text{Hash}(s, u)$  is uniform and independent of  $\text{ct}^*$  by the random-oracle property.  $\square$

### 8.2.5 Sizes and Comparison

Table 4 compares LSN-KEM with selected post-quantum KEMs.

Table 4: KEM size comparison.

| Scheme    | $ \text{pk} $ | $ \text{ct} $ | Assumption    | Failure      |
|-----------|---------------|---------------|---------------|--------------|
| Kyber-512 | 800 B         | 768 B         | Module-LWE    | $2^{-164}$   |
| HQC-128   | 2.25 KB       | 4.50 KB       | Syndrome Dec. | $2^{-128}$   |
| LSN-80    | 1.78 KB       | 288 B         | LSN           | $< 2^{-89}$  |
| LSN-128   | 2.78 KB       | 288 B         | LSN           | $< 2^{-157}$ |

### 8.2.6 CCA Security

We obtain IND-CCA security via the Fujisaki–Okamoto (FO) transform [FO99, HHK17] with implicit rejection. Let  $\text{BaseKEM} = (\text{KeyGen}, \text{Encaps}, \text{Decaps})$  be the LSN-KEM of section 8.2.2; recall that  $\text{Encaps}$  is deterministic once the random pair  $(s, u)$  is fixed. Define a hash function  $H : \{0, 1\}^* \rightarrow \{0, 1\}^\lambda$  (modelled as a random oracle).



**CCA-KeyGen**( $1^\lambda$ ):  $(\text{pk}, \text{sk}) \leftarrow \text{KeyGen}(1^\lambda)$ . Output  $\text{sk}' = (\text{sk}, d)$  where  $d \leftarrow \{0, 1\}^\lambda$  is an implicit-rejection secret.

**CCA-Encaps**( $\text{pk}$ ): Choose  $s \leftarrow \{0, 1\}^\lambda$  and  $u \leftarrow \{0, 1\}^K$ , compute  $(c_0, k_0) \leftarrow \text{Encaps}(\text{pk}; (s, u))$ , set  $c = c_0$  and  $K = H(s, u, c_0)$ . Output  $(c, K)$ .

**CCA-Decaps**( $\text{sk}', c$ ): Parse  $c = (s, \text{syn})$ . Recompute the blocks from  $s$  and  $v'_j = \text{majority}(\mathbf{1}_L(x))$  as in base decaps. Compute  $c' = \text{syn} \oplus v'$  and decode  $u' = \text{PolarSCLDecode}(c')$ . If  $u' = \perp$ , output  $K = H(d, c)$  (*implicit rejection*). Re-encapsulate  $(c'_0, k'_0) \leftarrow \text{Encaps}(\text{pk}; (s, u'))$ . If  $c'_0 = c$ , output  $K = H(s, u', c)$ ; otherwise output  $K = H(d, c)$  (*implicit rejection*).

**Theorem 8.3** (IND-CCA). *If BaseKEM is  $\delta$ -correct with  $\delta \leq 2^{-80}$  and IND-CPA secure, then CCA-KEM is IND-CCA secure in the random-oracle model.*

*Proof sketch.* The FO transform preserves IND-CPA security and adds CCA robustness via the re-encapsulation check. Implicit rejection ensures that decapsulation-failure events are indistinguishable from pseudorandom outputs; we inherit the multi-challenge ROM bound of [HHK17]. CCA security implies non-malleability, so tampered ciphertexts are rejected by the re-encapsulation check. The concatenated code affects only the correctness error  $\delta$ ; since  $\delta \leq 2^{-80}$  (Theorem 8.1), the standard FO analysis [HHK17] applies unchanged.  $\square$

### 8.2.7 Multi-User Security

In the multi-user setting,  $N$  parties generate independent key pairs  $(\text{pk}_i, \text{sk}_i)$  with independent secrets  $M_i \leftarrow \text{Sym}(n, \mathbb{F}_2)$ . The adversary receives all  $N$  public keys and may query decapsulation oracles for any user.

**Reduction to single-user security.** All users share only the public symplectic matrix  $S$ ; every other parameter ( $\text{seed}_i$ , noise  $e_{i,j}$ , permutation  $\pi_i$ ) is independent. Because the  $M_i$  are independent and the PRG seeds are independent, the  $N$  instances are *statistically independent* conditioned on  $S$ . By a standard hybrid argument over the  $N$  users, any adversary  $\mathcal{A}$  against  $N$ -user IND-CCA security with advantage  $\varepsilon_N$  yields an adversary against single-user IND-CCA security with advantage  $\varepsilon_1 \geq \varepsilon_N/N$ .

**Key-reuse analysis.** A more subtle concern is repeated encapsulation to the *same* user. Over  $q$  encapsulations the adversary sees  $q \cdot m = qNr$  samples  $(x_j, y_j)$  from the same secret  $M$ . The brute-force decoder must still search  $2^{n(n+1)/2}$  symmetric matrices, and the conditional SQ lower bound (Theorem 6.6) requires  $q_{\text{sq}} \geq 2^{2n-O(1)}$  queries. Translating sample count to query count via  $q_{\text{sq}} = qNr$ , the number of encapsulations needed for non-negligible advantage satisfies

$$q \gtrsim \frac{2^{2n-O(1)}}{Nr}.$$

(We treat each public sample as at most one SQ query; the R2a experiment supports a  $2^{2n}$  *sample* threshold.) For  $n = 65$ ,  $N = 2048$ ,  $r = 11$  we have  $Nr \approx 2^{14.5}$ , giving  $q \gtrsim 2^{114}$ —far beyond any realistic threat. We therefore treat multi-user and key-reuse security as *straightforward* consequences of the single-user bound, with no special vulnerability introduced by the LSN structure.

## 8.3 Implementation Security

**Constant-time by construction.** All core LSN operations are native  $\mathbb{F}_2$  operations—XOR, AND, and bit-permute—which are straightforward to implement in constant time on modern CPUs. There is no modular reduction with an odd modulus, no discrete Gaussian sampling, no rejection sampling, and no variable-time branching on secret data. This makes a constant-time implementation *straightforward* rather than requiring careful engineering as in lattice schemes.

Table 5: Implementation-security comparison: Kyber vs. LSN.

| Operation         | Kyber                           | LSN                                                           |
|-------------------|---------------------------------|---------------------------------------------------------------|
| Modular reduction | Variable-time ( $q = 3329$ )    | <b>None</b> ( $\mathbb{F}_2$ )                                |
| NTT butterfly     | Twiddle-dependent memory access | <b>Fixed XOR</b> pattern                                      |
| Noise sampling    | Discrete Gaussian (rejection)   | <b>Uniform</b> bits                                           |
| Decoder           | Noisy polynomial rounding       | <b>Majority</b> + fixed-tree SC                               |
| Secret-key format | $2n^2$ bits (general basis)     | $\mathbf{n}(\mathbf{n} + \mathbf{1})/2$ bits (symmetric $M$ ) |

**Side-channel surface.** The only non-linear operation on secret data is the majority vote in the inner repetition decoder. A popcount implementation is data-independent on all modern architectures (ARM NEON, AVX-512 VPOPCNTDQ). The outer polar SC decoder traverses a fixed butterfly tree of depth  $\log_2 N$ ; the memory access pattern depends only on  $N$ , not on the secret  $M$  or the noise. Consequently, timing attacks on the decoder require exploiting micro-architectural effects in the popcount instruction itself—far narrower than the modular-reduction and Gaussian-sampling side channels of lattice KEMs.

**Current status.** A Python prototype validates parameters and algorithms. A production constant-time Rust implementation (including constant-time popcount, bit-sliced symplectic transforms, and a constant-time SC decoder) is planned for full  $N = 2048$  validation and KAT vector generation.

## 9 Reduction Barriers and the Standard-Model Gap

Table 6 summarizes the status of natural reductions.

Table 6: Reduction landscape for LSN.

| Class                                                      | Status    | Reason                                                                                          |
|------------------------------------------------------------|-----------|-------------------------------------------------------------------------------------------------|
| Linear (real)                                              | BLOCKED   | real-linear: $\mathbb{E}[q]$ $L$ -independent;<br>$\mathbb{F}_2$ -linear: $\leq (1 - 2p)2^{-n}$ |
| Polynomial degree $< n$                                    | BLOCKED   | Error $\geq 2^{-n}$                                                                             |
| Linear, $\rho = 2n$ (full)                                 | BLOCKED   | Transport lemma; output Gram $\equiv 0$<br>(Theorem 9.1)                                        |
| Linear, $(3/2 + \epsilon)n \leq \rho < 2n$<br>(near-full)  | BLOCKED   | Gram rank $\leq 2c \leq (1 - \epsilon)n$<br>(Theorem 9.2)                                       |
| Linear, $\rho \leq (3/2 - \epsilon)n$ ,<br>$m = \Theta(n)$ | RULED OUT | LPQR26 App. D (all fixed linear<br>reductions)                                                  |
| Linear, $\rho \leq (3/2 - \epsilon)n$ ,<br>$m = \omega(n)$ | OPEN      | Entropy regime; LPQR26 bound<br>insufficient                                                    |
| LPN( $k = \Theta(n^2)$ )                                   | VACUOUS   | Proposition 9.3 + BKW cost<br>$2^{\Omega(n^2/\log n)}$                                          |
| Multi-sample or<br>beyond-poly                             | OPEN      | No proof or construction known                                                                  |

Lu *et al.* [LPQR26] prove that linear reductions cannot reduce symplectic LPN to LPN (their §2.4 and Appendix D): for any *fixed*  $B \in \mathbb{F}_2^{m \times 2n}$  the matrix  $BA$  has entropy at most  $(1 - d)mn$  for a constant  $d$ , so randomizing the public matrix forces a high-entropy  $B$ , which drives the output error weight past the Shannon-converse threshold. Per their Appendix D, the proof targets the primary cryptographic regime  $m = \Theta(n)$  (output rows  $m = (1 + \epsilon)n$ ); for  $m = \omega(n)$  it shows error weight larger than  $1/2 - \delta$  for any constant  $\delta$ , which they note is not sufficient to fully

rule out decodability, though they state they expect the barrier to extend to all  $m = \text{poly}(n)$ . Our Theorem 6.8 gives a complementary single-query barrier of  $(1 - 2p)2^{-n}$  at any noise rate.

### 9.1 Full-rank linear reductions are impossible

LPQR26 leave open the case  $m = \omega(n)$  (their bound shows error weight  $> 1/2 - \delta$ , which they note is insufficient to fully rule out decodability). We close the full-rank sub-case unconditionally, via a transport lemma that LPQR26 do not use.

Consider a linear reduction that maps a symplectic LPN instance  $(A, y)$  to an LPN instance  $(U, z)$  by outputting  $U = BA$  for a public matrix  $B \in \mathbb{F}_2^{m \times 2n}$ . The input columns satisfy  $S_A = A^\top \Omega A = 0$ .

**Theorem 9.1** (Transport lemma for full-rank linear reductions). *Let  $B \in \mathbb{F}_2^{m \times 2n}$  with  $\text{rank}(B) = 2n$  and  $m \geq 2n$ . Let  $B^+ \in \mathbb{F}_2^{2n \times m}$  be any left inverse ( $B^+ B = I_{2n}$ ) and define  $M := (B^+)^T \Omega B^+ \in \mathbb{F}_2^{m \times m}$ . Then for every isotropic-column matrix  $A$ ,*

$$(BA)^\top M (BA) = 0.$$

*A uniformly random matrix  $C \in \mathbb{F}_2^{m \times n}$  satisfies  $C^\top M C = 0$  with probability at most  $2^{-n^2/2 + O(n)}$ . Consequently, the output distribution of any fixed full-rank linear reduction is supported on a proper algebraic subset of  $\mathbb{F}_2^{m \times n}$  and has statistical distance  $1 - 2^{-\Theta(n^2)}$  from the uniform matrix distribution used in LPN.*

*Proof.* The identity follows from

$$\begin{aligned} (BA)^\top M (BA) &= A^\top B^\top (B^+)^T \Omega B^+ B A \\ &= A^\top (B^+ B)^\top \Omega (B^+ B) A = A^\top \Omega A = 0. \end{aligned}$$

For the probability bound,  $M$  has rank  $2n$  (because  $B^+$  has rank  $2n$  and  $\Omega$  is non-degenerate). Choose an invertible  $P \in \mathbb{F}_2^{m \times m}$  such that  $P^\top M P = \text{diag}(J, \dots, J, 0, \dots, 0)$  with  $n$  copies of  $J = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ . Writing  $P^{-1}C = [C_1; C_2]$  with  $C_1 \in \mathbb{F}_2^{2n \times n}$ , the condition  $C^\top M C = 0$  is equivalent to  $C_1^\top \text{diag}(J, \dots, J) C_1 = 0$ , i.e. the column space of  $C_1$  is totally isotropic in the  $2n$ -dimensional symplectic space  $(\mathbb{F}_2^{2n}, \text{diag}(J, \dots, J))$ . The number of  $n$ -dimensional totally isotropic subspaces of  $\mathbb{F}_2^{2n}$  is  $|\text{Lagr}(2n, \mathbb{F}_2)| = \prod_{i=1}^n (2^i + 1) = 2^{n(n+1)/2 + O(1)}$ , while the total number of  $n$ -dimensional subspaces is the Gaussian binomial  $\begin{bmatrix} 2n \\ n \end{bmatrix}_2 = 2^{n^2 + O(n)}$ . Hence a uniformly random  $C_1$  (equivalently, a uniformly random  $C$ ) yields a totally isotropic column space with probability  $2^{-n^2/2 + O(n)}$ .  $\square$

Theorem 9.1 blocks full-rank linear reductions at *every*  $m \geq 2n$  and *every* noise rate. The next theorem closes the near-full-rank stratum as well.

**Theorem 9.2** (Transport lemma for near-full-rank linear reductions). *Let  $B \in \mathbb{F}_2^{m \times 2n}$  with  $\text{rank}(B) = \rho = 2n - c$  where  $1 \leq c \leq n - 1$ . Let  $K = \ker B$  ( $\dim K = c$ ). Then there exists a matrix  $M \in \mathbb{F}_2^{m \times m}$  such that for every isotropic-column matrix  $A$ ,*

$$\text{rank}((BA)^\top M (BA)) \leq 2c - \text{rank}(\Omega|_K) \leq 2c.$$

*In particular, an adversarial reduction can choose  $K$  isotropic (possible for all  $c \leq n$ ), in which case the Gram rank of the output is at most  $2c$  deterministically. A uniformly random matrix  $C \in \mathbb{F}_2^{m \times n}$  satisfies  $\text{rank}(C^\top M C) \leq 2c$  with probability at most  $2^{-\Omega((n-2c)^2)}$ .*

*Proof.* Rank-factorize  $B = \tilde{B}P$  with  $\tilde{B} \in \mathbb{F}_2^{m \times \rho}$  full column rank and  $P \in \mathbb{F}_2^{\rho \times 2n}$  full row rank. Then  $(BA)^\top M (BA) = (PA)^\top (\tilde{B}^\top M \tilde{B})(PA)$ , and  $Q := \tilde{B}^\top M \tilde{B}$  ranges over all of  $\mathbb{F}_2^{\rho \times \rho}$  (lift:  $M := (\tilde{B}^+)^T Q \tilde{B}^+$  for a left inverse  $\tilde{B}^+$ ). The set of matrices  $N$  of the form  $P^\top Q P$  is therefore exactly the set of bilinear forms vanishing on  $K$  from both sides. Fix a basis adapted to  $K \oplus V$

with  $\dim V = 2n - c$ . In this basis the standard symplectic form  $\Omega$  becomes a block matrix with blocks  $\Omega_{KK}, \Omega_{KV}, \Omega_{VK}, \Omega_{VV}$ . Adding a form  $N$  that kills  $K$  on both sides changes only the  $VV$  block; the blocks  $\Omega_{KK}, \Omega_{KV}, \Omega_{VK}$  remain fixed. We therefore obtain a *minimal-rank completion* problem for the matrix

$$E = \begin{pmatrix} \Omega_{KK} & \Omega_{KV} \\ \Omega_{VK} & E_{VV} \end{pmatrix},$$

where  $E_{VV}$  is free. The classical formula for the minimal rank of such a completion (classical minimal-rank-completion folklore) gives

$$\min_{E_{VV}} \text{rank}(E) = \text{rank} \begin{pmatrix} \Omega_{KK} & \Omega_{KV} \end{pmatrix} + \text{rank} \begin{pmatrix} \Omega_{KK} \\ \Omega_{VK} \end{pmatrix} - \text{rank}(\Omega_{KK}).$$

Since  $\Omega$  is invertible, any  $c$  rows (resp. columns) of  $\Omega$  are linearly independent, so the two row- and column-ranks are exactly  $c$ . Hence

$$\min_{E_{VV}} \text{rank}(E) = 2c - \text{rank}(\Omega|_K).$$

The minimum is achieved constructively by  $E_{VV} = \Omega_{VK} \cdot G \cdot \Omega_{KV}$  with  $G$  any generalized inverse of  $\Omega_{KK}$  (over  $\mathbb{F}_2$ , obtained from a rank factorization). A symmetric g-inverse always exists:  $G^* := G^\top \Omega_{KK} G$  satisfies  $\Omega_{KK} G^* \Omega_{KK} = \Omega_{KK}$  and  $(G^*)^\top = G^*$ , yielding a symmetric  $E_{VV}$  and hence a symmetric matrix  $M$  representing the optimal form in the original coordinates. This justifies the corank-probability bound, which assumes a symmetric Gram matrix.

For the probability bound, a uniformly random  $n \times n$  symmetric matrix over  $\mathbb{F}_2$  has corank  $k$  with probability  $2^{-\Theta(k^2)}$ . Hence

$$\Pr[\text{rank}(C^\top M C) \leq 2c] \leq 2^{-\Omega((n-2c)^2)}.$$

This is  $1 - 2^{-\Omega(n^2)}$  only when  $n - 2c = \Omega(n)$ , i.e.  $\rho \geq (3/2 + \epsilon)n$  for any constant  $\epsilon > 0$ .  $\square$

**Rank stratification.** Together, Theorems 9.1 and 9.2 give a complete stratification of the fixed- $B$  linear-reduction landscape at any  $m$  and any noise rate:

- $\rho = 2n$  (full rank): **BLOCKED** — output Gram is identically zero (Theorem 9.1).
- $(3/2 + \epsilon)n \leq \rho < 2n$  (near-full): **BLOCKED** — output Gram rank  $\leq 2c \leq (1 - \epsilon)n$  (Theorem 9.2), while the uniform Gram has rank  $\approx n$  w.h.p.
- $\rho \leq (3/2 - \epsilon)n$  (mid/low rank): **OPEN** — the rank test dies ( $2c \geq (1 + \epsilon)n$ ). At  $m = \Theta(n)$ , LPQR26 rule out *all* fixed linear reductions (Appendix D, Theorem D.1); at  $m = \omega(n)$  their bound is insufficient. The entropy argument of Proposition 9.3 gives a complementary barrier. A quantitative bridge theorem that covers *all*  $B$ -distributions simultaneously would close this strip entirely.

This stratification applies to *fixed/public*  $B$  (per-realization). If the reduction draws  $B$  from a high-entropy distribution conditioned on  $BA$  being uniform, the matrix is randomized—but then the label noise  $Be$  is killed by piling-up at  $p = 1/4$  (row weight  $w$  gives bias  $2^{-w}$ ). The two regimes are endpoint observations; the quantitative trade-off over all  $B$ -distributions is an open problem (see Section 10).

A sharper, unconditional barrier comes from an information-theoretic floor:

**Proposition 9.3** (Entropy floor). *Any reduction from  $\text{LSN}(n)$  to  $\text{LPN}(k)$  that preserves the secret with recovery probability  $\alpha(n)$  must satisfy*

$$k \geq \alpha(n) H(L) - O(1) = \alpha(n) \Theta(n^2).$$

*In particular, constant success probability requires  $k = \Omega(n^2)$ , and success probability  $1 - o(1)$  requires  $k \geq H(L) - o(n^2)$ .*

*Proof.* The LSN secret space is  $\text{Lagr}(2n, \mathbb{F}_2)$  with  $|\text{Lagr}| = \prod_{i=1}^n (2^i + 1)$ . Hence  $H(L) = \log_2 |\text{Lagr}| = n(n+1)/2 + \Theta(1)$ , where the additive constant is  $\log_2 \prod_{i=1}^\infty (1 + 2^{-i}) \approx 1.25$ . If a reduction outputs an LPN( $k$ ) instance from which  $L$  is recovered with probability  $\alpha$ , Fano's inequality gives  $k \geq \alpha H(L) - H_2(1 - \alpha) = \alpha H(L) - O(1)$ .  $\square$

**Concrete-security win-win.** For near-perfect recovery of  $L$  (i.e.,  $\alpha = 1 - o(1)$  in Proposition 9.3), the entropy floor forces  $k \geq H(L) - o(n^2)$ . For the parameter sets in Table 1 this gives  $k \geq 863$  (80-bit),  $k \geq 2147$  (128-bit),  $k \geq 4755$  (192-bit), or  $k \geq 8387$  (256-bit). The best known attack on LPN( $k$ ) (BKW [BKW03]) requires time  $2^{\Omega(k/\log k)}$ . At these dimensions this yields  $2^{88.5}$ ,  $2^{194.0}$ ,  $2^{389.3}$ , and  $2^{643.5}$  respectively—all exceeding the target security levels. We stress that this is a concrete-security observation, not an asymptotic impossibility: BKW on LPN( $\Theta(n^2)$ ) runs in  $2^{\Theta(n^2/\log n)}$ , which is asymptotically smaller than brute-force on LSN ( $2^{\Theta(n^2)}$ ). For the parameter range of interest, however, the entropy floor lands BKW well above the security budget.

**Single-sample adaptivity does not increase degree.** A natural question is whether *adaptive* degree- $D$  feature-map reductions circumvent the static degree barrier. For the restricted model in which the reduction is *single-sample*—it may choose each  $\phi_j$  adaptively, but applies it only to the raw samples  $(a, b)$ , with the choice of  $\phi_j$  made as a function of previous answers—the answer is negative: in each round the feature map  $\phi_j(a, b)$  has degree at most  $D$ , and the effective degree of the entire transcript is bounded by  $D$ . When  $D < n$ , the Reed–Muller obstruction of ?? still applies: correlation with  $\mathbf{1}_L$  is at most  $2^{-n}$ , requiring exponentially many samples. Thus single-sample adaptive degree- $D$  feature-map reductions are blocked for every  $D < n$ , independent of the number of rounds.

The remaining open question is whether a *multi-sample* or *genuinely non-linear* adaptive reduction exists. Such a reduction could apply non-linear functions to *fresh* samples generated from previous outputs, or could use adaptively chosen non-linear queries in a way that raises the algebraic degree of the transcript. We do not claim impossibility for this richer model. We do note a win-win structure: a reduction in this richer model would likely improve random self-reductions for LPN itself, making it valuable even if it collapses LSN to the 6.5th family.

**The symplectic structure is reduction-inert.** A sharper observation comes from the structure of the LSN samples themselves. The deterministic quadratic constraint  $S_A = 0$  (the  $\Omega$ -Gram of the public matrix columns) is *public* and *x-free*: it depends only on the public matrix  $A$ , not on the secret  $x$  or the noise. Consequently,  $S_A = 0$  adds no  $x$ -information to an attacker and provides no usable handle for a reductionist trying to embed LPN into LSN. This mirrors the accepted status of Ring-LWE, which is regarded as a lattice-family assumption by *source* (ring structure native to the problem) rather than by reduction to plain LWE. We formalize this as:

**Conjecture 9.4** (Source novelty). *The following three invariants of LSN have no analogue in standard LPN, and each is destroyed by every natural reduction class that has been proposed:*

1. **Symplectic-Fourier self-duality.** *The indicator  $\mathbf{1}_L$  satisfies  $\mathcal{F}_\Omega[\mathbf{1}_L] = 2^n \mathbf{1}_L$ , where  $\mathcal{F}_\Omega$  is the symplectic Fourier transform. A standard linear function  $f(x) = \langle s, x \rangle$  has no such fixed-point property; its Fourier transform is a Dirac mass, not a function.*
2. **x-free quadratic constraint (sympLPN).** *In the sympLPN variant (Definition 4.3) the public matrix  $A$  satisfies  $S_A = 0$  (the  $\Omega$ -Gram of columns is identically zero). This is a deterministic, public, quadratic constraint on the public parameters alone. Our KEM uses pseudorandom rather than isotropic columns, so this invariant does not apply to the KEM instance distribution; it is a structural distinction of the sympLPN variant studied by LPQR26. LPN has no analogous public structure linking its samples.*

3. **Stabilizer degeneracy.** Every LSN secret  $L$  is a stabilizer state with stabilizer weight  $\Omega(n)$ . The secret space is the Lagrangian Grassmannian, not a vector space. LPN secrets are bare linear functions.

Each invariant is preserved by the symplectic group  $\text{Sp}(2n)$  and destroyed by any of the blocked reductions: linear maps lose the self-duality (Lu et al. [LPQR26]), degree- $D < n$  polynomial maps cannot represent  $\mathbf{1}_L$  (Reed–Muller distance), and entropy-compressing maps discard the stabilizer structure (Proposition 9.3).

Conjecture 9.4 is falsifiable: a single explicit reduction (membership-LSN, batch-LSN, or  $\text{sympLPN}$ )  $\rightarrow$  LPN that preserves any of the three invariants would refute the conjecture in its current form. It is not decidable by reduction analysis alone; it rests on the absence of such a reduction and the structural barriers documented above.

**The scrambling barrier at constant noise.** The LPQR26 linear barrier applies at any noise rate, and at our constant noise  $p = 1/4$  its error-amplification step is in fact *stronger*: a randomizing row of weight  $w$  leaves per-bit bias  $(1 - 2p)^w = 2^{-w}$  (piling-up), versus  $\approx 1 - 2wp$  in the low-noise regime. The residual gaps are therefore not about noise: (i) linear reductions consuming  $m = \omega(n)$  samples, where the LPQR bound does not fully rule out decodability (their Appendix D); and (ii) non-linear scramblers. For (ii) an algebraic obstruction remains:  $S_A = 0$  is a deterministic quadratic relation among the public vectors (our observation, distinct from the LPQR entropy argument). Direct inspection of the  $n = 3$  case shows that simple scramblers fall into one of three classes: a linear scrambler *transports* the quadratic relation ( $S'_A = 0$  with respect to the conjugated form  $M^{-\top} \Omega M^{-1}$ , so it remains detectable), a non-linear scrambler destroys the linear label structure required by LPN, and a random scrambler destroys the secret entirely. None of these simple classes simultaneously destroys  $S_A = 0$  and creates a fresh LPN secret-bit structure. The residual question — whether a sufficiently complex non-linear scrambler exists — is the same as Open Problem 2.

## 10 Open Problems

We state precisely the problems whose resolution would close the standard-model gap.

1. **Statistical dimension under  $S_A = 0$  for  $\text{sympLPN}$  [LPQR26].** The full SQ proof is blocked by the deterministic quadratic symplectic relation  $S_A = 0$  (the  $\Omega$ -Gram of columns) that characterises the  $\text{sympLPN}$  variant (Definition 4.3). Prove that conditioning on  $S_A = 0$  does not reduce the statistical dimension below  $2^{\Omega(n)}$ .
2. **LPN(low-noise)  $\rightarrow$  LSN.** In the cryptographically relevant regime  $p = O(1/\sqrt{n})$ , the reduction of [KLP+25] is vacuous. Does a non-vacuous reduction exist, or can its non-existence be proved?
3. **LWE  $\rightarrow$  LSN.** A reduction from LWE (or worst-case lattice problems) to LSN would provide the gold-standard foundation currently missing.
4. **Quantum lower bound.** Prove Conjecture 7.1: any quantum algorithm solving LSN requires  $\Omega(2^n)$  queries.
5. **Optimal polar-code rate.** Maximize the rate of the outer polar code while keeping decryption failure below  $2^{-128}$ .
6. **Pencil-extremality conjecture.** Prove Conjecture 6.5: that isotropic pencils ( $k \leq 2$ ) are extremal at scale  $|\text{Lagr}(2n, \mathbb{F}_2)|/2^{2n}$ , giving average correlation  $\leq 5\rho_{\text{avg}}$  for every subset of that size. This would upgrade Theorem 6.6 from a conditional to an unconditional worst-case bound.

7. **Sample freshness for rerandomized LSN.** Exact secret rerandomization is possible: applying a public symplectic matrix  $S \in \text{Sp}(2n)$  to the public part of each sample maps the secret  $L$  to  $S \cdot L$  while preserving labels. Does LSN admit a public transformation that yields *fresh* (statistically independent) samples for the rerandomized secret? A positive answer would yield a tight multi-user security bound and improve the reduction in Section 8.2.7. A negative answer would explain why the hybrid argument loses a factor  $N$ .

## 10.1 Honest Limitations of the Present Work

We flag five practical gaps that a reader should weigh against the theoretical claims above.

1. **Empirical gap at  $N = 2048$ .** The concatenated code is designed via the Bhattacharyya bound, which guarantees  $P_e \leq 2^{-80}$  for the target block length  $N = 2048$ . We have validated the SC decoder empirically for  $N \in \{128, 256, 512\}$  (all BLER = 0 in 200 trials), but a direct Monte-Carlo run at  $N = 2048$  is computationally expensive and has not been completed. The correctness claim for  $N = 2048$  therefore rests on the Bhattacharyya bound and the smaller-scale validation, not on a direct measurement.
2. **No constant-time reference implementation.** While section 8.3 argues that LSN-KEM is constant-time “by construction,” this property has not been verified on a concrete, audited implementation. A production-grade Rust reference implementation with known-answer test (KAT) vectors is planned but not yet available. Physical fault-injection resistance (e.g., bit-flip attacks on the majority vote) is likewise deferred to the audited implementation.
3. **Full-protocol SNARK is future work.** Section 8.1 optimizes only the *membership* sub-circuit ( $n^2$  constraints). A complete zero-knowledge proof of correct encapsulation—including the polar encoder, the majority-vote decoder, and the random permutation—remains open. The comparison in Table 2 should be read as a lower bound on what a full SNARK would cost.
4. **Loose multi-user reduction.** Section 8.2.7 uses a hybrid argument losing a factor  $N$ . A tight reduction would require sample freshness for a rerandomized secret. Exact secret rerandomization *is* possible—applying a public symplectic matrix  $S \in \text{Sp}(2n)$  to the public part of each sample maps  $L$  to  $S \cdot L$  while preserving labels—but *sample freshness* is blocked: because  $\mathbf{1}_L$  is non-linear, one cannot generate independent fresh samples for the new secret from the old public samples alone without knowing  $L$ . In practice the margin is comfortable ( $q \gtrsim 2^{114}$ ), but the bound is not tight.
5. **Practical optimizations deferred.** The canonical noise rate  $p = 1/4$  is fixed; whether LSN remains viable at other rates is unanalyzed. The inner repetition code uses a majority-vote decoder for constant-time implementability, but MAP decoding would achieve a lower effective noise. No public-key compression analogous to structured LPN is known. These are engineering tuning degrees of freedom, not structural vulnerabilities.

**Threshold cryptography.** The LSN secret  $M$  admits a natural linear secret-sharing scheme over  $\mathbb{F}_2$  ( $M = \sum M_j$ ), and the public samples  $(x_i, y_i)$  provide a publicly verifiable consistency check because the linear relation  $b = Ma = \sum_j M_j a$  is additively shareable. This enables distributed key generation where  $n$  parties jointly sample  $M$  without any single party learning it. Formalizing robustness against malicious minorities is future work.

**Quantum ECC interoperability.** Every LSN secret  $L$  is a stabilizer state ( $k = 0$ ). Its minimum stabilizer weight is  $\Omega(n)$  with high probability by the quantum Gilbert–Varshamov bound, so the secret carries natural quantum error-correcting structure. Whether this classical–quantum linkage yields concrete protocols (e.g., quantum-authenticated key exchange) is an open application direction.

## A Full Proof of Pairwise Correlation

Here we give a self-contained derivation of Lemma 6.1. Let  $D_0$  be the distribution on  $\mathbb{F}_2^{2n} \times \mathbb{F}_2$  in which  $a$  is uniform and  $b \sim \text{Bernoulli}(p)$  independently. For a fixed Lagrangian  $L$ , the LSN distribution  $D_L$  has density

$$D_L(a, b) = \frac{1}{2^{2n}} \left( (1-p) \mathbf{1}_{b=\mathbf{1}_L(a)} + p \mathbf{1}_{b \neq \mathbf{1}_L(a)} \right).$$

The likelihood ratio is

$$\frac{D_L(a, b)}{D_0(a, b)} = 1 + \mathbf{1}_L(a) \cdot \beta \cdot \left( \mathbf{1}_{b=1} - \mathbf{1}_{b=0} \frac{p}{1-p} \right),$$

where  $\beta = (1-2p)/p$ . Therefore

$$\ell_L(a, b) = \mathbf{1}_L(a) \cdot \beta \cdot \left( \mathbf{1}_{b=1} - \mathbf{1}_{b=0} \frac{p}{1-p} \right).$$

For two secrets  $L, L'$  with  $j = \dim(L \cap L')$ ,

$$\begin{aligned} \langle D_L, D_{L'} \rangle &= \mathbb{E}_{D_0}[\ell_L \ell_{L'}] \\ &= \mathbb{E}_a[\mathbf{1}_L(a) \mathbf{1}_{L'}(a)] \cdot \beta^2 \mathbb{E}_b \left[ \left( \mathbf{1}_{b=1} - \mathbf{1}_{b=0} \frac{p}{1-p} \right)^2 \right] \\ &= 2^{j-2n} \cdot \beta^2 \left( p + (1-p) \frac{p^2}{(1-p)^2} \right) \\ &= 2^{j-2n} \cdot \frac{(1-2p)^2}{p(1-p)}. \end{aligned}$$

Setting  $p = 1/4$  gives the coefficient  $(1/2)^2 / (1/4 \cdot 3/4) = (1/4) / (3/16) = 4/3$  exactly.

## B Q-Binomial Distance Distribution

The number of  $k$ -dimensional subspaces of  $\mathbb{F}_2^n$  is the Gaussian binomial coefficient

$$\begin{bmatrix} n \\ k \end{bmatrix}_2 = \prod_{i=0}^{k-1} \frac{2^{n-i} - 1}{2^{k-i} - 1}.$$

To count pairs of Lagrangians with intersection dimension  $k$ , fix a  $k$ -dimensional isotropic subspace  $W$ . The number of Lagrangians containing  $W$  equals the number of Lagrangians in the symplectic quotient  $W^\perp/W \cong \mathbb{F}_2^{2(n-k)}$ , namely  $\prod_{i=1}^{n-k} (2^i + 1)$ . Summing over all  $k$ -dimensional subspaces  $W \subseteq L$  and using the identity  $\prod_{i=1}^{n-k} (2^i + 1) = 2^{(n-k)(n-k+1)/2} \prod_{i=1}^{n-k} (1 + 2^{-i})$ , the exact probability is

$$\Pr[j = k] = \frac{\begin{bmatrix} n \\ k \end{bmatrix}_2 \cdot 2^{(n-k)(n-k+1)/2}}{\prod_{i=1}^n (2^i + 1)}.$$

The moment generating function  $C_n = \mathbb{E}[2^j]$  converges rapidly to 2 as  $n \rightarrow \infty$ .



## C Reduction Barrier Summary

**Linear.** For any linear classifier  $h(a) = \langle w, a \rangle + t$ , the advantage

$$\text{Adv}(h) = \left| \Pr_{D_L}[h(a) = 1] - \Pr_{D_0}[h(a) = 1] \right|$$

is zero because  $\Pr_{D_L}[h(a) = 1] = \Pr_{D_0}[h(a) = 1]$ .

**Polynomial degree  $< n$ .** The exact polynomial representation of  $\mathbf{1}_L$  has degree exactly  $n$  and  $2^n$  monomials. Any polynomial of degree  $d < n$  cannot vanish on  $L^c$ , so its correlation with  $\mathbf{1}_L$  is at most  $2^{-n}$ .

**Adaptive linear SQ.** For real-linear queries  $q(a, b) = \langle w, a \rangle + cb$  we have  $\mathbb{E}_{D_L}[q] = c(p + 2^{-n}(1 - 2p))$ , independent of  $L$ . For  $\mathbb{F}_2$ -linear queries the maximum advantage is  $(1 - 2p)2^{-n}$  (Theorem 6.8). In either case the advantage is exponentially small in  $n$ , so adaptive linear SQ algorithms require exponentially many samples.

## D Extension to $\mathbb{F}_q$

All results in the main text are stated over  $\mathbb{F}_2$ . Here we sketch the extension to arbitrary finite fields  $\mathbb{F}_q$ , which widens the parameter space and allows a trade-off between block length  $n$  and field size  $q$ .

### D.1 Symplectic Spaces over $\mathbb{F}_q$

Let  $q$  be a prime power and  $V = \mathbb{F}_q^{2n}$  equipped with a non-degenerate alternating bilinear form  $\Omega(x, y) = x^T J y$  where  $J = \begin{pmatrix} 0 & I_n \\ -I_n & 0 \end{pmatrix}$ . A subspace  $L \subset V$  is *Lagrangian* if it is maximal isotropic; then  $\dim L = n$  and  $|\text{Lagr}(2n, \mathbb{F}_q)| = \prod_{i=1}^n (q^i + 1)$ .

For fixed  $L$  and  $j = \dim(L \cap L')$ , the number of Lagrangians  $L'$  with intersection dimension  $j$  is

$$N_j(n, q) = \begin{bmatrix} n \\ j \end{bmatrix}_q \cdot q^{(n-j)(n-j+1)/2},$$

where  $\begin{bmatrix} n \\ j \end{bmatrix}_q$  is the Gaussian binomial coefficient. Summing over  $j$  recovers  $|\text{Lagr}(2n, \mathbb{F}_q)|$ .

### D.2 Exact Correlation over $\mathbb{F}_q$

The label remains binary ( $b \in \mathbb{F}_2$ ) even though the ambient space is  $\mathbb{F}_q^{2n}$ , so the SQ machinery of Feldman *et al.* applies without modification. The likelihood ratio calculation is identical with 2 replaced by  $q$ :

**Lemma D.1** ( $\mathbb{F}_q$  pairwise correlation). *For  $L, L' \in \text{Lagr}(2n, \mathbb{F}_q)$  with  $j = \dim(L \cap L')$ ,*

$$\langle D_L, D_{L'} \rangle = \frac{(1 - 2p)^2}{p(1 - p)} \cdot \frac{q^j}{q^{2n}}.$$

**Lemma D.2** ( $\mathbb{F}_q$  average correlation). *Let  $C_{n,q} = \mathbb{E}[q^j]$  computed via the  $q$ -binomial distribution. Then*

$$\rho_{\text{avg}} = \frac{(1 - 2p)^2}{p(1 - p)} \cdot C_{n,q} \cdot q^{-2n}.$$

*Numerically,  $C_{n,q} \rightarrow 2$  as  $n \rightarrow \infty$  for every  $q$  (the limit is independent of the field size).*

### D.3 Security Parameters over $\mathbb{F}_q$

The SQ lower bound becomes  $q_{\text{queries}} \geq q^{2n-O(1)}$ . One “bit” of security corresponds to  $\log_2(q)$  field operations, so the effective bit-security is  $2n \log_2(q)$ . Table 7 gives the block lengths  $n$  needed for target security levels.

Table 7: LSN parameters over  $\mathbb{F}_q$  for  $p = 1/4$ .

| Security | $\mathbb{F}_2$ : $n$ | $\mathbb{F}_3$ : $n$ | $\mathbb{F}_5$ : $n$ | $\mathbb{F}_7$ : $n$ |
|----------|----------------------|----------------------|----------------------|----------------------|
| 80-bit   | 41                   | 26                   | 18                   | 15                   |
| 128-bit  | 65                   | 41                   | 28                   | 24                   |
| 192-bit  | 97                   | 62                   | 42                   | 35                   |
| 256-bit  | 129                  | 82                   | 56                   | 46                   |

Larger  $q$  reduces the required  $n$  but increases the cost of each sample (field arithmetic in  $\mathbb{F}_q$  versus  $\mathbb{F}_2$ ). The optimal choice depends on the target platform.

### D.4 Open Questions

1. Extend to  $q$ -ary labels  $y \in \mathbb{F}_q$  with noise over  $\mathbb{F}_q$ ; this requires  $q$ -ary SQ lower bounds.
2. The proof holds for  $q = p^k$  with  $k > 1$  because symplectic theory is field-agnostic.
3. In characteristic 2, alternating is equivalent to symmetric; in odd characteristic they differ. The intersection distribution formula remains valid in all characteristics.

## References

- [AAC+22] C. Adjiman, N. Aragon, S. Bettaieb, L. Bidoux, O. Blazy, J.-C. Deneuville, P. Gaborit, S. Ghosh, A. Hauteville, and G. Zemor. HQC. NIST PQC Round 4 Submission, 2022.
- [ABD+21] M. Albrecht, D. Bernstein, T. Chou, C. Cid, J. Gilcher, T. Lange, V. Maram, I. von Maurich, R. Misoczki, R. Niederhagen, K. Paterson, E. Persichetti, C. Peters, P. Schwabe, N. Sendrier, J. Szefer, C. Tjhai, M. Tomlinson, and W. Wang. Classic McEliece. NIST PQC Round 4 Submission, 2021.
- [ABW10] A. Akavia, S. Goldwasser, and V. Vaikuntanathan. Simultaneous hardcore bits and cryptography against memory attacks. In *TCC*, pages 474–495. Springer, 2010.
- [Ale11] A. Aleknovich. More on average case vs approximation complexity. *Computational Complexity*, 20(4):755–786, 2011.
- [Ari09] E. Arikan. Channel polarization: A method for constructing capacity-achieving codes for symmetric binary-input memoryless channels. *IEEE Trans. Inform. Theory*, 55(7):3051–3073, 2009.
- [BBB+18] B. Bünz, J. Bootle, D. Boneh, A. Poelstra, P. Wuille, and G. Maxwell. Bulletproofs: Short proofs for confidential transactions and more. In *IEEE S&P*, pages 315–334, 2018.
- [BCTV14] E. Ben-Sasson, A. Chiesa, D. Tromer, and M. Virza. Scalable zero knowledge via cycles of elliptic curves. In *CRYPTO*, pages 276–294. Springer, 2014.
- [BFJ+94] A. Blum, M. Furst, J. Jackson, M. Kearns, Y. Mansour, and S. Rudich. Weakly learning DNF and characterizing statistical query learning using Fourier analysis. In *STOC*, pages 253–262. ACM, 1994.

- [BFKL94] A. Blum, A. Furst, M. Kearns, and R. Lipton. Cryptographic primitives based on hard learning problems. In *CRYPTO*, pages 278–291. Springer, 1994.
- [BKW03] A. Blum, A. Kalai, and H. Wasserman. Noise-tolerant learning, the parity problem, and the statistical query model. *J. ACM*, 50(4):506–519, 2003.
- [BPR12] A. Banerjee, C. Peikert, and A. Rosen. Pseudorandom functions and lattices. In *EUROCRYPT*, pages 719–737. Springer, 2012.
- [BCGI18] E. Boyle, G. Conteau, N. Gilboa, and Y. Ishai. Compressing vector OLE. In *ACM CCS*, pages 896–912, 2018.
- [CRSS98] A. R. Calderbank, E. M. Rains, P. W. Shor, and N. J. A. Sloane. Quantum error correction via codes over  $\text{GF}(4)$ . *IEEE Trans. Inform. Theory*, 44(4):1369–1387, 1998.
- [DKS17] I. Diakonikolas, D. M. Kane, and A. Stewart. Statistical query lower bounds for robust estimation of high-dimensional Gaussians and Gaussian mixtures. In *FOCS*, pages 73–84. IEEE, 2017.
- [DGM14] N. Döttling, J. Müller-Quade, and A. Nascimento. IND-CCA secure cryptography based on a variant of the LPN problem. In *ASIACRYPT*, pages 485–508. Springer, 2014.
- [DP12] Y. Dodis, K. Pietrzak, and D. Wichs. Key derivation without entropy waste. In *EUROCRYPT*, pages 93–110. Springer, 2012.
- [EKM17] A. Esser, R. Kübler, and A. May. LPN decoded. In *CRYPTO*, pages 486–514. Springer, 2017.
- [FO99] E. Fujisaki and T. Okamoto. Secure integration of asymmetric and symmetric encryption schemes. In *CRYPTO*, pages 537–554. Springer, 1999.
- [FGR+17] V. Feldman, E. Grigorescu, L. Reyzin, and S. Vempala. Statistical query lower bounds for robust estimation of high-dimensional Gaussians and Gaussian mixtures. In *FOCS*, pages 211–220. IEEE, 2017.
- [FV19] A. Fazeli and A. Vardy. On the scaling exponent of binary polarization kernels. In *IEEE ISIT*, pages 1472–1476, 2019.
- [GGPR13] R. Gennaro, C. Gentry, B. Parno, and M. Raykova. Quadratic span programs and succinct NIZKs without PCPs. In *EUROCRYPT*, pages 626–645. Springer, 2013.
- [GJL14] Q. Guo, T. Johansson, and C. Löndahl. Solving LPN using covering codes. In *ASIACRYPT*, pages 1–20. Springer, 2014.
- [GL22] S. Gollakota and D. Liang. PAC-learning noisy stabilizer states in the statistical query model. *arXiv preprint*, 2022.
- [Got97] D. Gottesman. Stabilizer codes and quantum error correction. PhD thesis, Caltech, 1997.
- [Gro96] L. K. Grover. A fast quantum mechanical algorithm for database search. In *STOC*, pages 212–219. ACM, 1996.
- [Gro16] J. Groth. On the size of pairing-based non-interactive arguments. In *EUROCRYPT*, pages 305–326. Springer, 2016.
- [HHK17] D. Hofheinz, K. Hövelmanns, and E. Kiltz. A modular analysis of the Fujisaki–Okamoto transformation. In *TCC*, pages 341–371. Springer, 2017.

- [HKL+12] S. Heyse, E. Kiltz, V. Lyubashevsky, C. Paar, and K. Pietrzak. Lapin: An efficient authentication protocol based on ring-LPN. In *FSE*, pages 346–365. Springer, 2012.
- [HVU14] S. H. Hassani, R. Urbanke, and S. V. Macris. Near-optimal finite-length scaling for polar codes over large alphabets. *IEEE Trans. Inform. Theory*, 60(11):6960–6976, 2014.
- [Kea98] M. Kearns. Efficient noise-tolerant learning from statistical queries. *J. ACM*, 45(6):983–1006, 1998.
- [Kir11] P. Kirchner. Improved generalized birthday attack. *Cryptology ePrint Archive*, Report 2011/377, 2011.
- [KLP+25] A. Khesin, J. Lu, A. Poremba, S. Ramkumar, and V. Vaikuntanathan. Hardness of decoding random quantum stabilizer codes. *arXiv:2509.20697*, 2025.
- [LF06] E. Levieil and P.-A. Fouque. An improved LPN algorithm. In *SCN*, pages 348–359. Springer, 2006.
- [LPQR26] J. Lu, A. Poremba, Y. Quek, and S. Ramkumar. Post-quantum cryptography from quantum stabilizer decoding. *arXiv:2603.19110*, 2026.
- [Lyu05] V. Lyubashevsky. The parity problem in the presence of noise, decoding random linear codes, and the subset sum problem. In *APPROX-RANDOM*, pages 378–389. Springer, 2005.
- [Mac71] I. G. Macdonald. Symmetric Functions and Hall Polynomials. Oxford University Press, 1971.
- [McE78] R. J. McEliece. A public-key cryptosystem based on algebraic coding theory. *DSN Progress Report*, 44:114–116, 1978.
- [MT09] R. Mori and T. Tanaka. Performance and construction of polar codes on symmetric binary-input memoryless channels. In *IEEE ISIT*, pages 1496–1500, 2009.
- [NC00] M. A. Nielsen and I. L. Chuang. Quantum Computation and Quantum Information. Cambridge University Press, 2000.
- [NIS22] NIST. Post-quantum cryptography standardization, 2022. <https://csrc.nist.gov/projects/post-quantum-cryptography>.
- [PHGR13] B. Parno, J. Howell, C. Gentry, and M. Raykova. Pinocchio: Nearly practical verifiable computation. In *IEEE S&P*, pages 238–252, 2013.
- [PQS26] A. Poremba, Y. Quek, and P. Shor. Learning stabilizers with noise. *arXiv:2410.18953*, 2025.
- [Pra62] E. Prange. The use of information sets in decoding cyclic codes. *IRE Trans. Inform. Theory*, 8(5):5–9, 1962.
- [Reg05] O. Regev. On lattices, learning with errors, random linear codes, and cryptography. *J. ACM*, 56(6):34, 2005.
- [SAB+22] P. Schwabe, R. Avanzi, J. Bos, L. Ducas, E. Kiltz, T. Lepoint, V. Lyubashevsky, J. M. Schanck, G. Seiler, and D. Stehlé. CRYSTALS-Kyber. NIST PQC Round 3 Submission, 2022.
- [Sei15] M. Seidl. Polar coding: Finite-length aspects. PhD thesis, Friedrich-Alexander-Universität Erlangen-Nürnberg, 2015.

- [Sho94] P. W. Shor. Algorithms for quantum computation: Discrete logarithms and factoring. In *FOCS*, pages 124–134. IEEE, 1994.
- [Tri12] P. Trifonov. Efficient design and decoding of polar codes. *IEEE Trans. Commun.*, 60(11):3221–3227, 2012.
- [TV15] I. Tal and A. Vardy. List decoding of polar codes. *IEEE Trans. Inform. Theory*, 61(5):2213–2226, 2015.
- [WS21] A. Wang and F. Zhang. Revisiting the BKW algorithm for LPN. *Cryptology ePrint Archive*, Report 2021/213, 2021.
- [ZJW16] B.-Y. Zhang, L. Jia, and M.-W. Wang. Security analysis of LPN under perfect code cover. *IEEE Access*, 4:3979–3986, 2016.
- [komm] A. Krohn. komm: An open-source library for communication systems. <https://github.com/ryukau/komm>, 2024.